

OpenCPI

HDL Development Guide

Revision History

Revision	Description of Change	Date
1.01	Creation from content in previous OpenCPI Component Development document	2016-02-1
1.02	Add new content to be consistent with current system as of February 2016	2016-02-23
1.03	Processed review comments for 3 reviewers	2016-02-24
1.10	Missing details about environment, time service, dependencies	2016-04-15
1.2	2017Q1 reviews, changed simulator model	2017-02-24
1.3	Update for 2017.Q2 issues	2017-08-08
1.4	Updates for 2018.Q1: primitive library clarifications, more backpressure explanation	2018-02-26
1.5	Update for device2device container connections, other clarifications	2018-09-24

Table of Contents

1	References.....	6
2	Overview.....	7
3	HDL Workers.....	10
3.1	Execution Model for HDL Workers.....	10
3.2	Creating an HDL Worker.....	11
3.3	Building the Worker.....	12
3.4	Worker Makefile.....	13
3.5	HDL Worker Description File: the OWD XML File.....	14
3.5.1	HDL Worker OWD Top Level Attributes.....	14
3.5.2	Attributes of Worker Interfaces.....	16
3.5.3	Clocking Elements and Attributes.....	17
3.6	The Authored Worker: the VHDL Architecture or Verilog Module Body.....	18
3.6.1	Signal Naming Conventions and Data Types.....	19
3.7	The Control Interface to the HDL Worker.....	21
3.7.1	XML attributes for the HDL Worker's Control Interface.....	21
3.7.2	Reset Behavior and Initializations in HDL Workers and Infrastructure.....	22
3.7.3	Clock and Reset Input Signals in the Control Interface.....	22
3.7.4	Life cycle/Control Operation Signals in the Control Interface.....	22
3.7.5	Property Access.....	27
3.7.6	Property Data Types.....	29
3.7.7	Raw Access to Properties.....	30
3.7.8	Summary of HDL Worker Control Interface.....	32
3.8	HDL Worker Data Interfaces for OCS Data Ports.....	33
3.8.1	Message Payloads vs. Physical Data Width on Data Interfaces.....	34
3.8.2	Byte Enables on Data Interfaces.....	35
3.8.3	Streaming Data Interfaces to/from the HDL Worker.....	36
3.8.4	Message Data Interfaces.....	46
3.8.5	HDL Worker Data Interface Summary.....	46
3.9	Time Service Interface.....	48
3.10	Memory Service Interfaces.....	49
4	Building HDL Assets.....	50
4.1	HDL Build Targets.....	50
4.2	The HDL Build Hierarchy.....	52
4.3	HDL Directory Structure.....	56
4.4	HDL Search Paths when Building.....	58
4.4.1	Searching for HDL Primitives.....	58
4.4.2	Searching for XML files (OCS, OPS) when Building Workers.....	59
4.4.3	Searching for Workers in Component Libraries.....	60
5	HDL Primitives.....	62
5.1	HDL Primitive Libraries.....	63
5.1.1	Source Files in Primitive Libraries.....	63
5.1.2	Package Declarations for Primitive Libraries.....	64
5.1.3	Instantiating Modules in Primitive Libraries.....	65

5.1.4	Providing Target-specific or Vendor-specific Versions of Primitive Modules.....	65
5.1.5	Exporting and Using the Results of Building HDL Primitive Libraries.....	67
5.2	HDL Primitive Cores.....	68
6	HDL Assemblies for Creating Bitstreams/Executables.....	70
6.1	The HDL Assembly XML file.....	72
6.2	The Makefile for Building an Assembly.....	75
6.3	Specifying the Containers that Implement the Assembly on Platforms.....	76
6.4	HDL Container XML files.....	78
6.4.1	Service Connections in a Container.....	79
6.4.2	Preparing the Bitstream/Executable Artifact File.....	80
7	HDL Simulation Platforms.....	82
7.1	Execution of Simulation Bitstreams and Containers.....	83
8	HDL Device Naming.....	85
8.1	PCI-based HDL devices.....	86
8.2	Ethernet-based HDL Devices.....	87
8.3	Simulator Device Naming.....	88
9	The ocpihdl Command Line Utility for HDL Development.....	89
9.1	General, non-Worker Commands for the ocpihdl Utility.....	91
9.1.1	admin command – Print the Device’s Administrative Information.....	91
9.1.2	bram command – Create a Configuration BRAM File from an XML File.....	91
9.1.3	deltatime command – Perform Time Synchronization Test on Device.....	91
9.1.4	emulate command – Emulate a Network-based Device (admin space only).....	91
9.1.5	ethers command – Display Available (Up and Connected) Network Interfaces.....	91
9.1.6	probe command – Test Existence and Availability of Device.....	92
9.1.7	load command – Load Bitstream.....	92
9.1.8	getxml command – Retrieve XML Metadata from Device.....	92
9.1.9	radmin command – Read a Specific Address in Device’s Admin Space.....	92
9.1.10	reset command – Perform Soft Reset on Device.....	92
9.1.11	rmeta command – Read from Addresses in the Metadata Space of the Device.....	92
9.1.12	search command – Search for all Available HDL Devices.....	92
9.1.13	unbram command – Create an XML File from a Config BRAM File.....	92
9.1.14	wadmin command – Write Specific Addresses in the Device’s Admin Space.....	93
9.1.15	settime command – Set Device’s Time from System Time.....	93
9.2	Worker Commands: Commands that Operate on Individual Workers.....	94
9.2.1	get command – Get All or Single Worker Instance Property Information.....	94
9.2.2	set command – Set Property Value in a Worker Instance.....	94
9.2.3	control command – Change Control State of Worker Instance.....	94
9.2.4	status command – Display the Status of a Worker Instance.....	94
9.2.5	Primitive worker commands using worker indices.....	95
9.2.6	wclear command – Clear a Worker’s Error Registers.....	95
9.2.7	wdump command – Dump a Worker’s Status Registers (not Properties).....	95
9.2.8	wop command – Perform a Control Operation on a Worker.....	95
9.2.9	wread command – Read a Worker’s Configuration Property Space.....	95
9.2.10	wreset command – Assert Reset for a Worker.....	96
9.2.11	wunreset command – Deassert Reset for a Worker.....	96
9.2.12	wwctl command – Write a Worker’s Control Register.....	96

9.2.13	wwpage command – Write a Worker’s Window/Page Register.....	96
9.2.14	wwrite command – Write a Worker’s Configuration Property Space.....	96
10	HDL Platform and Device Development.....	97

1 References

This document depends on several others. Primarily, it depends on the ***OpenCPI Component Development Guide (CDG)***, which describes concepts and definitions common to all OpenCPI authoring models.

Table 1: References to Related Documents

Title	Published By	Link
OpenCPI Overview	OpenCPI	Public URL: https://github.com/opencpi/opencpi/raw/2018.Q3/doc/pdf/OpenCPI_Overview.pdf
OpenCPI Component Development Guide	OpenCPI	Public URL: https://github.com/opencpi/opencpi/raw/2018.Q3/doc/pdf/OpenCPI_Component_Development.pdf
OpenCPI RCC Development Guide	OpenCPI	Public URL: https://github.com/opencpi/opencpi/raw/2018.Q3/doc/pdf/OpenCPI_RCC_Development.pdf

2 Overview

This document describes how to develop component implementations, known as **HDL workers**, for FPGAs, using the **Hardware Description Language (HDL) Authoring Model**. Workers are typically created in a component library, so that they are available for OpenCPI application developers and users.

Some knowledge of FPGA terminology is assumed here, but this document is also useful for non-FPGA developers in understanding the OpenCPI FPGA development process.

This document builds on the information provided in the **OpenCPI Component Development Guide (CDG)**, which introduces concepts and processes used for OpenCPI component development in general.

In addition to describing how to develop **HDL workers**, this document also describes:

- how to create **HDL primitive libraries**, which are libraries of smaller/simpler reusable code modules sometimes used in the design of HDL workers
- how to create **HDL primitive cores**, which are prebuilt and possibly pre-synthesized modules sometimes incorporated into HDL workers
- how to assemble a group of connected HDL workers to form an **HDL assembly**, which is realized in a complete FPGA bitstream.

HDL assemblies enable an FPGA to execute a subset of, or all of, the components specified in an OpenCPI application. Usually when an OpenCPI application uses an FPGA, it is using it to execute some of the components specified in the application, with the others executing on other processors typically using other authoring models.

However, in some cases an HDL assembly provides workers for *all* the components required by an application.

For HDL development, the **OpenCPI Component Development Kit (CDK)** utilizes technology-specific FPGA synthesis and simulation tools (e.g. Xilinx ISE and Isim, Altera Quartus, Modelsim etc.).

The following sections describe the development of HDL workers, primitives (libraries and cores), assemblies and bitstreams to support the execution of parts of component-based applications on FPGAs. All these terms are prefixed with HDL here to avoid confusion when they are used elsewhere.

HDL Authoring Model: the OpenCPI authoring model targeting Hardware Description Languages that are appropriate for FPGA development, currently using VHDL, with some legacy support for Verilog. New HDL workers should be written in VHDL. Full support for Verilog and System-Verilog is not currently supported.

HDL Target: a particular type of FPGA device, usually what is considered a part family, that is the target of compilation or synthesis, where the result can be used for any architecturally similar device. Examples are “virtex6”, or “stratix4”, or “zynq”. Simulators are also HDL targets, including Mentor’s Modelsim and Xilinx Isim.

HDL Worker: an HDL (e.g. VHDL) implementation of a component specification, with the source code written according to HDL authoring model. While most HDL workers are *application* workers (usable in portable applications), a special type is *device* workers which are for controlling hardware physically attached to the FPGA. For application workers, it is common and recommended to have an RCC worker that also implements the same spec. Such workers are sometimes called “work-alikes”.

HDL Primitive: an HDL asset that is lower level than workers, that is used as a building block for workers. HDL primitives can either be *libraries* or *cores*.

HDL Primitive Library: a collection of low level modules compiled from source code that can be referenced in HDL worker code. An HDL worker declares which HDL primitive libraries it draws modules from.

HDL Primitive Core: a low level module that may be built and/or synthesized from source code, or imported as presynthesized and possibly encrypted from 3rd parties, or generated by tools like Xilinx CoreGen or Altera MegaWizard. An HDL worker declares which primitive cores it requires (and instantiates).

The following diagram shows the hierarchy of modules when an FPGA design is realized using OpenCPI:

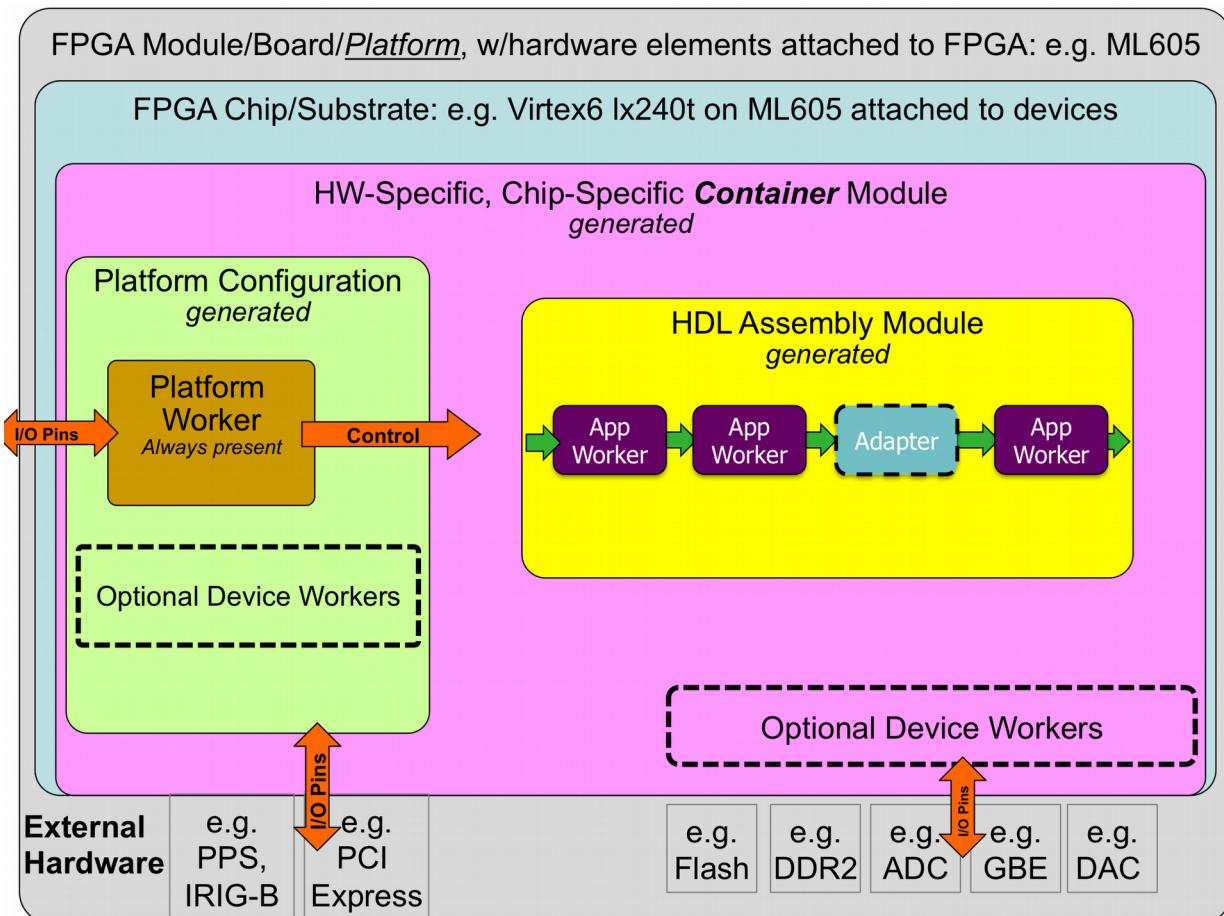


Figure 1: HDL Hierarchy

HDL Assembly: a composition of connected HDL workers that are built into a complete FPGA configuration **bitstream**, acting as an OpenCPI artifact. The resulting bitstream is executed on an FPGA to implement some part or all of (the components of) the overall OpenCPI application. The HDL code is *automatically generated* from a brief description in XML.

HDL Platform: an OpenCPI platform based on an FPGA that is enabled to host OpenCPI HDL workers. Simulators are also considered HDL platforms.

HDL Platform Worker: a specific type of HDL worker providing infrastructure for implementing control/data interfaces to devices and interconnects external to the FPGA or simulator (e.g. PCI Express, Clocks). See the [**HDL Platform and Device Development**](#) section.

HDL Device Worker: a specific type of HDL worker that supports external devices attached to FPGAs. See the [**HDL Platform and Device Development**](#) section.

HDL Platform Configuration: a prebuilt (presynthesized) assembly of device-level HDL workers that represent a particular configuration of device support modules for a given HDL platform. The HDL code is *automatically generated* from a brief description in XML. See the [**HDL Platform and Device Development**](#) section.

HDL Container: a complete design for an entire FPGA, which includes an HDL assembly and an HDL platform configuration combined in a specified fashion. The HDL code is *automatically generated* from a brief description in XML.

HDL development in OpenCPI includes both *application workers* in a component library, which perform functions independent of any specific hardware attached to the FPGA, as well as *device workers* that are designed to support specific external hardware such as ADCs, flash memories, I/O devices, etc.

HDL device workers are developed as part of enabling an HDL platform (an FPGA on a particular board) for OpenCPI. See the [**HDL Platform and Device Development**](#) section.

The sections below for HDL/FPGA development are:

- Developing application workers in a component library
- The HDL Build Process: building HDL assets for different target devices and platforms
- The HDL Build Hierarchy: how whole device “bitstreams” are created from other assets
- Developing assemblies of workers on FPGAs.

3 HDL Workers

This section describes how to write an HDL worker, and defines what distinguishes HDL worker development from developing workers using other authoring models. It builds on the worker development section in the Component Development Guide (CDG) that describes what is common to building workers for all authoring models.

HDL workers can consist of a single standalone source code module or reference and instantiate lower level models: HDL primitive libraries and HDL primitive cores. In either case a worker is compiled and (when appropriate) synthesized for a given **HDL Target**, as described in detail in the [HDL Build Targets](#) section. HDL worker source code cannot reference other workers.

3.1 Execution Model for HDL Workers

An HDL worker executes when enabled from its control interface. When not in the operating control state a worker should only execute in response to control operations. When in the operating state, it is expected to operate continuously until a control operation changes its state (via **stop** or **release** control operation). As a convenience, the code generated shell has its own implementation of the start control operation and sets the **is_operating** signal accordingly. It is strongly recommended that every worker use that signal to enable its operation. Execution of workers normally processes data arriving at input ports and produces data at output ports, possibly changing the values of volatile properties.

All workers are provided a reset signal, which is guaranteed to be asserted for at least 16 control clock cycles. If a worker needs more time or clock cycles to perform its initialization, then it must implement the **initialize** control operation as described in [Lifecycle/Control Operation Signals](#).

All HDL workers are provided the control clock which is used to control the operating state of the work, access property registers, and may be used for data processing. However, this may restrict performance in some cases since the control clock may not be the fastest clock rate that the logic can support.

Workers may be written with different clocks used at different ports, allowing some ports to operate at clock rates different from the control clock. When a worker is written this way it is up the author to perform appropriate clock domain synchronization inside the worker between control signals using the control clock and signals operating using different clocks. In many cases the only signals that need to “cross the clock domain boundary” inside such a worker are the control **reset** and **is_operating** signals.
[This multiple-clock support is preliminary in the current framework].

3.2 Creating an HDL Worker

The process of writing a new HDL worker (after the OCS exists), starts with using the `ocpidev create worker` command, as described in the CDG, to create the worker's directory and its initial content, usually as a subdirectory of a component library. This command can be executed in a project's directory, a component library's directory or even a completely separate directory not part of any pre-existing library or project. The name of an HDL worker always has the `.hdl` suffix, and the language must be specified as **VHDL** or **Verilog**. Languages are case insensitive. The default is **vhdl**. Here are some examples of using `ocpidev` to create HDL workers:

```
ocpidev create worker xyz.hdl -L vhdl
ocpidev -v -l dsplib create worker
ocpidev -S fft2d-spec create worker fft.hdl
```

As described in the CDG, this will create an initial worker description XML file (OWD), `<worker>.xml`, in the worker directory, which is subsequently edited for worker-specific attributes. Similarly, a **Makefile** is also created in the worker's directory, which may also be edited to enter worker-specific non-default build options. Creating an HDL worker also generates various files in the worker's **gen** subdirectory, including the code skeleton file which is initially copied to the `<worker>. <language-suffix>` in the worker's directory. The first `ocpidev` command above would result in the following directory tree for the worker:

```
xyz.hdl/
    Makefile          # Makefile for this worker
    xyz.xml           # OWD for this worker
    xyz.vhd           # editable source code for this worker
    gen/xyz-skel.vhd # initial skeleton for this worker
        xyz-defs.vhd # definitions enabling instantiation
        xyz-impl.vhd # the generated shell of the worker
        xyz.build      # the generated build configuration file
```

None of the files in the **gen** subdirectory should be edited. Since this directory and its contents are all generated by tools, all are deleted when the `make clean` command is issued here or at the project or library level. For HDL workers using VHDL, the initial worker code file (`xyz.vhd`), that you must subsequently edit, contains only the **architecture** of the worker. The **entity** declaration of the worker was automatically generated for you, and is found in the generated file `gen/xyz-impl.vhd`.

3.3 Building the Worker

The above command establishes the worker's directory (`xyz.hdl`), its `Makefile`, its OWD file (`xyz.xml`), and its initial source code file (`xyz.vhd`). The worker is built by issuing the

```
make <targets>
```

command in the worker's directory, or

```
make <targets> xyz.hdl
```

in the library directory containing the worker's directory, assuming the worker was created in a library. In both the above cases `<targets>` would specify which HDL platforms or part families to build for. If a default for `<targets>` was already specified in the project's `Project.mk` file, nothing would be necessary here.

These commands will compile (or synthesize) the worker for the active set of **HDL Targets**, such as `zynq`, `virtex6`, `stratix4`, `isim` or `modelsim`. HDL targets are described in detail in the [Building HDL Assets](#) section.

The worker may be built as it was originally created by `ocpidev`, for any targets, before adding any code to implement the actual function of the worker.

As with any type of worker, compilation output is placed in the `target-TTT` subdirectory of the worker, for each target in the currently active set. An example explicitly specifying the targets is:

```
make Hd1Targets="modelsim virtex6"
```

This would compile (and for non-simulation targets, synthesize) the worker for two targets. At this point, after building the initial generated worker code file, the directory representing the new worker looks like this:

```
xyz.hdl/
    Makefile          # Makefile for this worker
    xyz.xml          # OWD for this worker
    xyz.vhd          # editable source code for this worker
    gen/xyz-skel.vhd # initial skeleton for this worker
        xyz-defs.vhd # definitions enabling instantiation
        xyz-impl.vhd # the generated shell of the worker
        xyz.build      # the generated build config file
    target-virtex6/
        xyz.ngc        # virtex6 synthesized core file
        xyz-xst.out    # log of tool output
    target-modelsim/
        xyz/*          # modelsim compilation result files
        xyz-modelsim.out # log of tool output
```

The `gen` directory and all the `target-*` subdirectories are generated, and should not be edited. More details about these files are described below. As they are generated files, they are removed when the `make clean` command is issued.

3.4 Worker `Makefile`

The worker `Makefile` is usually not hand-edited, but there are cases where editing is necessary. Several variables are available for all authoring models, as described in the CDG. If your HDL worker will consist of multiple source files, you can add VHDL or Verilog source files (beyond the main file whose name is the worker name) by specifying the `SourceFiles` variable in the `Makefile`. Even if the primary language of the worker is VHDL, other Verilog files can be used, and vice versa. In addition to the `make` variables mentioned in the CDG for all authoring models (e.g. `SourceFiles`, `Libraries`, `XmlIncludeDirs`, `OnlyTargets`, `ExcludeTargets`), the HDL worker `Makefile` may also include these additional ones:

Table 2: HDL Worker `Makefile` Variables

Variable Name in HDL Worker <code>Makefile</code>	Usable as default in library <code>Makefile</code> ?	Description
<code>Cores</code>	N	A list of HDL primitive cores built elsewhere
<code>VerilogIncludeDirs</code>	Y	Searchable directories for Verilog include files, in addition to the worker directory
<code>HdlExactPart</code>	N	A variable to override the default part within a family specified by <code>HdlTarget</code> (s). See <i>HDL Build Targets</i> .

For HDL workers, the `Libraries` variable described in the CDG refers to HDL primitive libraries, described in the [*HDL Primitive Libraries*](#) section.

When using the `Libraries` and `Cores` variables, if a name in the list has no slashes, it is found by searching the path for HDL primitives as described in the [*HDL Search Paths*](#) section. If the name *does* contain slashes, it is the specific path name of the library or core, usually a relative path name within the same project.

3.5 HDL Worker Description File: the OWD XML File

This file specifies characteristics for an HDL worker which expand on those found in the spec file (OCS). In many cases it can be empty or left as it is. A default version is generated when the worker is created. Aspects of the OWD that are common to all authoring models are described in the CDG. Aspects specific to the HDL authoring model are described here.

The primary reasons to customize this OWD XML file for an HDL worker are:

- **Add implementation-specific properties**

You can add additional worker properties for the implementation, beyond what is in the spec file. The **property** element accomplishes this, with the same xml format as in the OCS file.

- **Add more accessibility to OCS properties**

You can add more access capabilities for existing properties, via the **specproperty** element. E.g. make a property that is write-only in the OCS to be also readable in the implementation for debug purposes. The only attributes are **name** and the access properties: **readable**, **writable**, **volatile**, **initial**.

- **Specify that an OCS property is in fact a parameter (compile-time) property**

You can indicate that in this worker, the OCS property is actually a compile-time value. This only applies to **initial** properties in the OCS. When this is specified, you can also specify the default value of the parameter property.

- **Specify which control operations this worker will implement**

You can specify which control operations are in the implementation: none are required. This uses the **ControlOperations** attribute of the OWD.

- **Specify interface style and implementation attributes for data ports**

You can specify whether the port uses a stream or message interface, and provide additional details for those interfaces (e.g. data path width, or whether the port supports aborting messages). See [Attributes of HDL Worker Data Ports](#).

Only the last item above is in fact specific to HDL workers. All the rest apply to workers of all authoring models and are fully described in the CDG.

3.5.1 HDL Worker OWD Top Level Attributes

The top level of the HDL OWD is the XML **HdlWorker** element, which can have the XML attributes in the table below (beyond those defined for all OWDs in the CDG). All are optional, and are only specified when the default behavior must be overridden. These attributes are in addition to those that are valid for all authoring models as described in the OWD XML description in the CDG.

Table 3: HdIWorker Attributes

HdIWorker Attribute Name	Data Type	Description
DataWidth	unsigned	The default physical width of data ports for this worker. Any individual port can override this. The default value when this attribute is not specified is based on the protocol (messages) defined for the port in the OCS.
RawProperties	boolean	A boolean value indicating whether the worker will use the raw property interface for all properties. The default is false. The raw interface is described below under Raw access to properties and is typically only used for device workers.
FirstRawProperty	string	A string value indicating the name of the first property that requires the raw property interface. Properties before this use the normal property interface.

The attributes below are for HDL infrastructure workers coded to the “outer” or OCP interfaces and not supported for general users.

Outer	boolean	Whether the worker implements the outer interface, used in internal OpenCPI modules or for legacy code.
Pattern	string	An external signal naming pattern (described below this table) for all signals of the worker. The default is "%s_ ", which indicates a prefix of the port name followed by underscore. External signals are those defined using the OCP interface standard, not the inner worker signals.
PortPattern	string	A port naming pattern used when port names and signal (not data) direction are used in the generated code. I.e. for each worker port, a naming pattern is defined both for input signals and output signals of the port. The default is "%s_%n", which indicates a prefix of the port name followed by underscore, and then in or out for signals that are input to the worker or output from the worker.
SizeOfConfig-Space	unsigned 64-bit	Overrides the size of the configuration space in bytes. The default is based on the actual properties.
Sub32BitConfig-Properties	boolean	Whether this worker needs to address items smaller than 32 bits (and thus requires byte enables in its interface). The default is based on the actual properties.

The **Pattern** and **PortPattern** attributes are like sprintf format strings in C, with **%s** being the port name (**%S** capitalized), **%n** being **in** or **out** for signal, not data, direction (**in** to or out of the worker), **%m** being **m** or **s** for master/slave (and **%M** for **M/S**), **%0** or **%1** for port ordinal within OCP profile (0 origin or 1 origin), **%i/I** for using **i/I** for input, **o/o**

for output, `%N` for `In` or `Out`, and `%w/w` for profile name (lower case or capitalized). These patterns are only used when exporting or importing OCP-based workers.

3.5.2 Attributes of Worker Interfaces

Other aspects of the OWD for HDL workers are described in the sections describing each interface: the [control interfaces](#), [data interfaces](#) (stream and message*), and [service interfaces](#) (time and memory*). They are summarized here, and described in detail later.

Table 4: HDL Worker Interface Attributes

Attribute Name	Which Interfaces?	Description
<code>Timeout</code>	Only control	Control clock cycles allowed for control ops
<code>MyClock*</code>	All but control	Declare that interface has its own clock
<code>Clock*</code>	All but control	Declare that interface uses a particular clock
<code>DataValueSize</code>	Any data	The minimum unit of data, in bytes.
<code>DataValueGranularity</code>	Any data	The minimum multiple of data values
<code>NumberOfOpCodes</code>	Any data	Maximum number of opcodes to support
<code>MaxMessageValues</code>	Any data	Maximum message length in “data values”
<code>ZeroLengthMessages</code>	Any data	Declare support for zero length messages
<code>Abortable</code>	Stream data	Declare that messages can be aborted.
<code>PreciseBurst</code>	Stream data	Declare that messages will have known length at the start of the message.
<code>SecondsWidth</code>	Time	Bits in the <code>seconds</code> field of time-of-day
<code>FractionWidth</code>	Time	Bits in the <code>fraction</code> field of time-of-day.
<code>AllowUnavailable</code>	Time	Declare tolerance for time to be unavailable

* Support for memory interfaces is preliminary/partial in the current framework.

3.5.3 Clocking Elements and Attributes

[Support for using clocks separate from the control clock is preliminary in the current framework]

The default clocking configuration of HDL workers is that all interfaces use the clock that is provided in the control interface. The control interface clock is always present.

All other interfaces can have a **MyClock** boolean attribute set to true, which indicates that the interface has its own clock which has no relationship with the control clock. I.e., its clock is in a different clock domain. In this case a **clk** input signal is included in its interface.

Interfaces may also have a **Clock** string attribute which identifies a specific clock to use for this interface, by name, which may be shared with other interfaces. The string value of this attribute is either the name of an OCS data port or the name of a separately defined clock (described just below). In the former case, the meaning is: the clock for this interface should be the same as the clock used at the other, identified interface.

A clock may be defined separate from any interface, and then referred to using the **clock** attribute of those interfaces. A separately defined clock may also be used independent of any interface. These separately defined clocks are needed when:

- A clock does not really belong to any interface, but may still be used by them.
- A clock has non-default attributes.

Separate clocks are defined with the **Clock** XML element, as a child element of the top level **HdLWorker** element. The attributes of Clock elements are:

Table 5: HDL Worker Clock Element Attributes

Attribute Name	Value Type	Description
Name	string	The name used to identify this clock in the clock attribute of other interfaces. Also used in forming the input signal name for this clock into this worker.
Signal	string	The actual signal name that should be used for this clock. The default is the clock name, in the name attribute.

3.6 The Authored Worker: the VHDL Architecture or Verilog Module Body

This section describes how and where to write the actual VHDL or Verilog source code for a worker.

The functional code or “business logic” of the worker is in the **architecture** section (VHDL), in the **xyz.vhd** file (and possibly subsidiary source files). In Verilog, it is in the body of the **module**, and the file is **xyz.v**. This architecture/module is initially generated as a skeleton of the **inner** worker. It is surrounded by an automatically generated logic **shell** which provides robust and composable interfaces compliant with the Open Core Protocol (OCP) interfaces defined for the entire **outer** worker.

This shell and the **entity** declaration for the inner worker are found in the generated file: **gen/xyz-impl.vhd** for VHDL and **gen/xyz-impl.vh** for Verilog. The skeleton file, consisting of an empty inner worker, becomes the authored worker when the functional logic is written/inserted into that file.

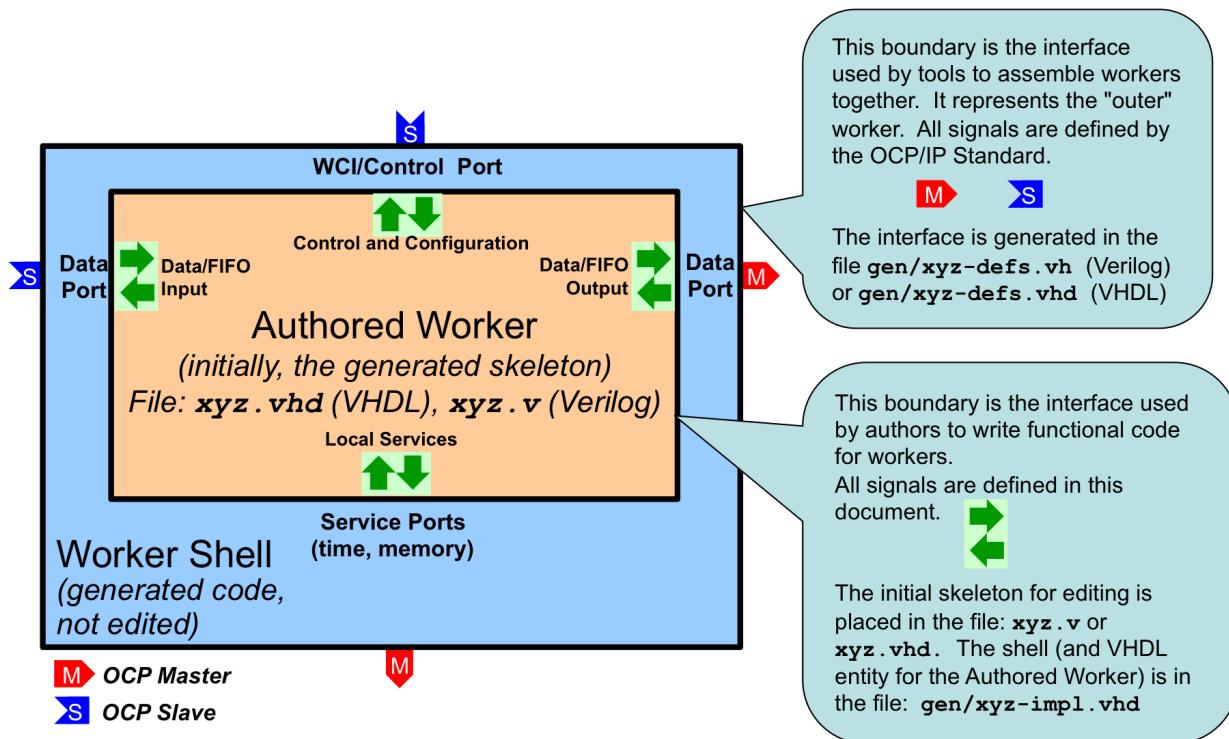


Figure 2: Worker Code and Files

All OpenCPI HDL workers are characterized by their properties, their data and service ports and their clocks, and usually the clocks are simply associated with ports, or even more simply, a single clock is commonly used with all ports. The job of implementing the inner worker is the job of:

- processing the various data ports’ inputs to the worker to
- produce the various data ports’ outputs of the worker,

- using the services provided at the service ports.

For each port of the worker (including the control port) there are input signals (into the worker) and output signals (out of the worker).

In VHDL, these groups of inputs and outputs are in a record type. Thus for each port (whether control, data producing, data consuming or other service) there is an ***input signal*** record and an ***output signal*** record named <port>_in and <port>_out, respectively. In Verilog there are no record types, so individual signals simply have the <port>_in_ and <port>_out_ prefixes.

E.g., with a “filter” worker that has a “sensor” input port, and a “result” output port, the VHDL entity declaration (in the `gen/filter-impl.vhd` file) would be:

```
entity filter_worker is
  port(
    ctl_in          : in  worker_ctl_in_t;
    ctl_out         : out worker_ctl_out_t;
    sensor_in       : in  worker_sensor_in_t;
    sensor_out      : out worker_sensor_out_t;
    result_in       : in  worker_result_in_t;
    result_out      : out worker_result_out_t);
end entity filter_worker;
```

The actual individual signals in each record depend on the contents of the OCS and OWD files. These signals will be described below. Note that the name of the “control port” defaults to `ctl`. An example skeleton file for this worker, in the file `filter.vhd`, would be:

```
library ieee; use ieee.std_logic_1164.all, ieee.numeric_std.all;
library ocpi; use ocpi.types.all;
architecture rtl of filter_worker is
begin
  -- put the logic for this worker here
end rtl;
```

Note that while the overall worker has the name “filter”, the entity being implemented in the architecture here is `filter_worker`, the *inner* worker.

The clause `use ocpi.types.all` introduces all the data types in that package to the namespace for the architecture code of the worker. This is most convenient for using the built-in types provided by OpenCPI. However, if the author wants to avoid any collisions with their own types or functions, they can remove this `use` clause and fully qualify references to types provided by OpenCPI.

3.6.1 Signal Naming Conventions and Data Types

Other than the property access signals described below, the signals in these worker interfaces are mostly a combination of IEEE `std_logic_vector` and a boolean type, `bool_t`, that is used for various boolean indicator signals. All these VHDL types and related constants are defined in the `ocpi.types` package from the `ocpi` HDL primitive library.

The VHDL type `bool_t` acts as much like the VHDL type `BOOLEAN` as possible (with various operator overloading functions), while still being based on `std_logic`. The `to_boolean` and `to_bool` functions explicitly convert to and from the VHDL `BOOLEAN` type, respectively. The `its` function is a convenient synonym for the `to_boolean` function, enabling code like:

```
if its(ready) then
  ...
end if;
```

There are also two constants for this type, `btrue` and `bfalse`. These types and constants may also be used in user-written primitives, and are used in code automatically generated by OpenCPI.

In VHDL, all signals into and out of the authored worker are in the `in` and `out` records of each port.

All data types created by OpenCPI use the `_t` suffix. All enumeration values defined by OpenCPI use the `_e` suffix.

OpenCPI uses the term **port** to mean a high level data flow interface in and out of all types of workers. This conflicts with the use of the term in VHDL and Verilog, which means the individual signals (of any type) that are the inputs and outputs of an entity (VHDL) or module (Verilog).

In this section on HDL workers, this document uses the term **interface** to be the HDL worker's set of input and output port signals that correspond to the high level OpenCPI ports as defined in the OCS and OWD for the HDL worker. We also use the term **interface** for the implicit control port of all workers. An HDL worker has a control interface (for the implicit control port), data interfaces (for the explicit data ports defined in the OCS), and service interfaces (for service ports as defined in the HDL worker's OWD).

3.7 The Control Interface to the HDL Worker.

Every HDL worker has a control interface that at a minimum provides a *control* clock and associated reset into the worker. Normally, the control interface also is used to:

- Convey life cycle *control operations* like `initialize`, `start` and `stop`
- Access the worker's *configuration properties* as specified in the OCS and OWD

In VHDL, when the default name of the control interface is used (`ctl`), the input signals are prefixed with `ctl_in`. and the output signals are prefixed with `ctl_out`. I.e. the input signals are in the `ctl_in` record port, and the output signals are in the `ctl_out` record port.

When the spec for the HDL worker (in its OCS) has the (rarely used) `NoControl` attribute set to true, only the clock and reset signals are present. In this case no signals associated with control operations or properties are present and there are no `ctl_in` or `ctl_out` signals. Only `wci_clk` for the clock signal and `wci_reset` for the reset signal are present.

3.7.1 XML attributes for the HDL Worker's Control Interface

Most aspects of the control interface are generically specified either in the OCS (e.g. the `NoControl` attribute), or at the top level of the OWD XML (e.g. the `ControlOperations` attribute). Several additional control interface attributes for HDL workers may be specified in a `ControlInterface` child element of the OWD. One example is the `Timeout` attribute described below. An example of an HDL OWD with this attribute would be:

```
<HdlWorker>
  <ControlInterface Timeout='100' />
  ...
</HdlWorker>
```

The following table contains attributes that may be specified for the `ControlInterface` element:

Table 6: HDL Worker ControlInterface Element Attributes

Attribute Name	Value Type	Description
<code>Timeout</code>	ULong	<p>The minimum number of control clock cycles that should be allowed for the worker to complete control operations or property access operations. When this number is exceeded the worker is considered inoperative and a timeout error is reported. The worker completes these operations using the <code>ctl_out.done</code> or <code>ctl_out.error</code> signals described below.</p> <p>The default value is 16.</p>

3.7.2 Reset Behavior and Initializations in HDL Workers and Infrastructure

Per (at least) Xilinx recommended practice, OpenCPI uses and generates resets synchronously,. Thus resets are implemented active high and are asserted and deasserted synchronously. There are several reasons for this policy, but one is that less logic is typically needed to implement the resetting of register state.

At power up or reconfiguration, resets are asserted, so they will be asserted on the *first* clock edge. Per the OCP specification, resets will always be asserted for at least 16 clock cycles.

If registers (state) truly need an initial value (e.g. for simulation cleanliness or glitch-free initialization or sim-vs-synth consistency), it is preferred to set an initial *default expression* value in VHDL or Verilog, rather than using asynchronous reset. This is done by providing an initial value expression in the signal declaration. Note that current Xilinx (ISE, and ISIM), Altera (Quartus), and Mentor (Modelsim) support such initialization without using any resources.

In the OpenCPI HDL infrastructure, applying resets to register state is only used to serve a functional purpose, and not the default practice.

3.7.3 Clock and Reset Input Signals in the Control Interface

The signal `ctl_in.clk` is the clock for all other control port signals as well as the *default* clock for all other data or service ports of the worker. The `ctl_in.reset` signal (asserted high) is asserted and deasserted synchronously with this clock. This `reset` is guaranteed to be asserted for 16 clock cycles. When 16 clocks are not enough to perform initialization, the worker should implement the `initialize` control operation (see below). The control reset, like all other resets generated by the OpenCPI infrastructure, is initially asserted.

If the worker (in its OWD) declares that other data or service ports have clocks that are different than this control clock (i.e. those interfaces operate in *different clock domains*), the worker implementation code has responsibility for the appropriate synchronizations between this control clock (and its associated signals) and any other signals related to the data or service interfaces. In particular, it is the worker's responsibility to propagate this control reset to the reset outputs associated with other interfaces, in their clock domain.

3.7.4 Life cycle/Control Operation Signals in the Control Interface

Other than the control `reset` signal, the life cycle of all workers is managed by life cycle control operations, according to the diagram below. When a worker's control reset is deasserted, it enters the `exists` state. Control operations cause state changes as shown. When control operations fail, the `unusable` state is entered. The worker autonomously enters the `finished` state, without any control operation.

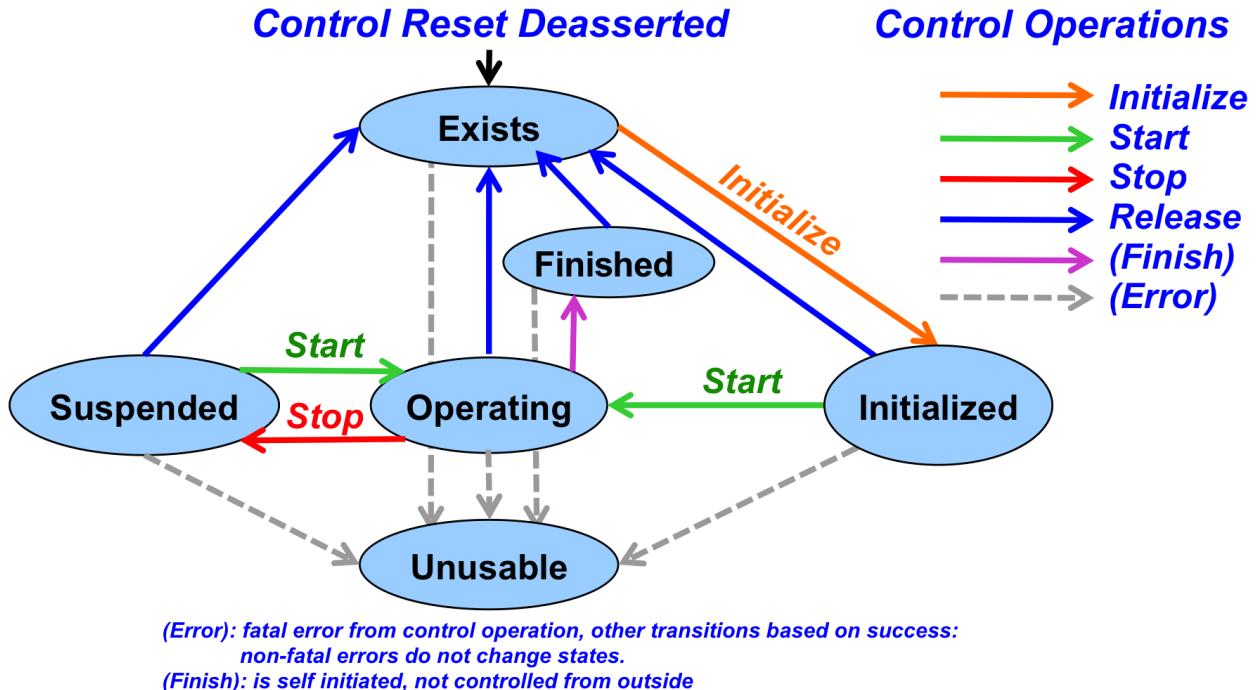


Figure 3: Control Operations and States

In simple and common cases, when the worker has no need to implement any of these operations, there is a single input signal indicating when the worker should *operate*, called `is_operating`. I.e., after `reset` is deasserted, the worker should *operate* only when this `is_operating` signal is asserted. Many HDL workers use this signal (as well as `clk` and `reset`) and no others in the control interface.

The `is_operating` signal indicates that the worker has been *started* and should now perform its function. A worker *must not* perform any data transactions at its data ports unless `is_operating` is true. This is necessary for robust system-level application control to suspend and resume all or parts of an application.

An example of a simple HDL worker would be:

```

architecture rtl of filter_worker is
  signal mystate_r : std_logic_vector(7 downto 0); -- some state
begin
  process (ctl_in.clk) is
  begin
    if rising_edge(ctl_in.clk) then
      if its(ctl_in.reset) then
        mystate <= "01010101";
      elsif its(ctl_in.is_operating) then
        -- do the clocked functions of this worker
      end if;
    end if;
  end process;
end rtl;
  
```

This allows the worker to be suspended and resumed since nothing happens during suspension, when `is_operating` is false.

If the `initialize` control operation is not implemented, then when the `reset` signal is deasserted, the worker is considered to be *initialized*. If `initialize` is implemented, the worker is considered in the `exists` state after `reset` is deasserted. The worker's OWD specifies whether `initialize` is implemented by this worker.

When a worker needs to explicitly support other control operations, there are two input and two output signals it may use. The `control_op` signal is a VHDL enumeration value that conveys which control operation is in progress. When there is no operation in progress, it has a value of `NO_OP_e`. Otherwise the choices are: `INITIALIZE_e`, `START_e`, `STOP_e`, `RELEASE_e`, `BEFORE_QUERY_e`, and `AFTER_CONFIG_e`. The operation is terminated by the worker asserting the `done` or `error` output signals, after which the control operation is considered accomplished successfully (if `done`) or not (if `error`). Note that the `done` signal is driven to a default value of `btrue` in the entity declaration and the `error` signal defaults to `bfalse`. The worker does not need to drive these at all if it will always perform the control operations in a single cycle and will never need to assert `error`. The operation will be forced to complete with a timeout error if neither `done` nor `error` is asserted within the number of control clock cycles indicated in the `Timeout` attribute of the `ControlInterface` element.

All VHDL types specifically associated with the control interface are in the `wci` package of the `ocpi` library, including the enumeration values just mentioned. E.g.:

```
library ocpi;
...
if ctl_in.control_op = ocpi.wci.START_e ...
```

A common example of a control operation might be when the worker needs multiple clock cycles to accomplish something like `initialize` or `start`. In that case it notices when the `control_op` signal changes from `no_op_e`, and then performs the operation, asserting `done` (or `error`) when the operation has completed. An example where initialization takes 10 clock cycles, would be:

```

architecture rtl of example_worker is
    signal init_count_r : unsigned(4 downto 0);
begin
    process (ctl_in.clk) is
    begin
        if rising_edge(ctl_in.clk) then
            if its(ctl_in.reset) then
                init_count <= (others => '0');
            elsif ctl_in.control_op = INITIALIZE_e then
                init_count <= init_count + 1;
            elsif its(ctl_in.is_operating) then
                -- do normal functions
            end if;
        end if;
    end process;
    -- initialize takes 10 clocks, all others take 1
    ctl_out.done <=
        '0' when ctl_in.control_op = INITIALIZE_e and init_count < 20
        else '1';
end rtl;

```

Another convenience input signal, `state`, indicates which life cycle state the worker is in. It changes when control operations succeed. It is a VHDL enumeration value: `EXISTS_e`, `INITIALIZED_e`, `OPERATING_e`, `SUSPENDED_e`, `FINISHED_e`, and `UNUSABLE_e`. These types are also in the `ocpi.wci` package.

Finally, there are two control output signals that the worker can use to indicate two other conditions. The first control output signal is `finished`. The worker uses this to indicate it has entered the `finished` state, and will perform no further work. This enables the worker to tell control software that its work is finished and perhaps that the application the worker is part of can be considered finished. This signal should be deasserted upon `reset`. Asserting `finished` will cause `is_operating` to become false, and `state` to become `FINISHED_e`.

The second, `attention`, allows the worker to indicate an interrupt or other condition to control software. This signal is for legacy compatibility and should not be used in new workers. It should be deasserted on `reset`.

Here is a summary of the control interface signals in the `ctl_in` record. The `bool_t` type is in the `ocpi.types` package, and the `control_op_t` and `state_t` are in `ocpi.wci`

Table 7: Control Input Signals

Signal	Type	Description
<code>clk</code>	<code>std_logic</code>	The clock for the control interface and the default clock for all other interfaces and ports.
<code>reset</code>	<code>bool_t</code>	Asserted high and synchronously, for the control interface, for at least 16 clocks. Initially asserted.
<code>control_op</code>	<code>control_op_t</code>	An enumeration type specifying the currently active control operation, with the value <code>no_op_e</code> when there is no active control operation. Control operations persist until <code>done</code> or <code>error</code> signal in the <code>ctl_out</code> record is true.
<code>state</code>	<code>state_t</code>	An enumeration type indicating the worker's current control state. Changes when control operation ends (via <code>done</code> or <code>error</code>) or <code>finished</code> is asserted.
<code>is_operating</code>	<code>bool_t</code>	Indicates the worker is started and is in an operating state. Persists until <code>stop</code> or <code>release</code> operation completes or <code>finished</code> or <code>reset</code> is asserted.
<code>abort_control_op</code>	<code>bool_t</code>	A command indicating that a long-duration control operation is being forcibly aborted. A pulse.
<code>is_big_endian</code>	<code>bool_t</code>	For dynamic endian workers, set at reset.

Here is a summary of the control interface signals in the `ctl_out` record:

Table 8: Control Output Signals

Signal	Type	Description
<code>done</code>	<code>bool_t</code>	Indicates the successful end of a control operation. The default value is true indicating that all control operations complete in the same cycle they start.
<code>error</code>	<code>bool_t</code>	The signal indicating the unsuccessful end of a control operation. Default value is false.
<code>finished</code>	<code>bool_t</code>	A persistent indication, not deasserted after being asserted, until reset, that the worker has entered the <code>finished</code> state. Default is false.

These signals are assigned the default value in the entity port declaration.

3.7.5 Property Access

A worker's configuration properties are accessed via two additional record port signals called `props_in` and `props_out` (separate from the `ctl_in` and `ctl_out` records for the control interface). The individual signals within these records depend on what types of properties have been declared in the OCS and OWD.

When the worker shell is generated based on the OCS and OWD, the accessibility of properties determines which registers and signals are generated and made available to the worker code for each property. If any property is specified as being a **raw** property in the OWD, then the raw interface is also generated, and used for all such properties, as described below in [raw access to properties](#). Otherwise the following rules apply:

Table 9: HDL Worker Property Logic Rules

Property is writable or initial	Property is Readable	Property is Volatile	Logic Description
Yes	No	No	Value is registered in the shell and register outputs are available in <code>props_in</code>
Yes	Yes	No	Value is registered in the shell and register outputs are available in <code>props_in</code> . The readback value is from the register outputs.
Yes	No	Yes	Value is registered in the shell and register outputs are available in <code>props_in</code> . The readback value is from the worker in <code>props_out</code> .
No	Yes	No	The readback value is from the worker in <code>props_out</code> , but is cached by control software since the worker is not expected to change it after it is operating.
No	No	Yes	The readback value is from the worker in <code>props_out</code> .

In the tables below, for a property called `foo`, the signals will be present as described. The types of the signals are all in the `ocpi.types` package.

The signals possibly present in the `props_in` record are in the following table.

Table 10: HDL Worker Property Input Signals

Signal in <code>props_in</code>	Included when:	Type	Signal Description
<code>foo</code>	Writable or Initial	*	The registered value last written by control software. The type is dependent on the property type.
<code>foo_length</code>	Writable or Initial and type is sequence	<code>ulong_t</code>	The registered 32 bit unsigned number of elements in the sequence when property is a sequence.
<code>foo_written</code>	Writable	<code>bool_t</code>	Indicates the entire value is being written. Persists until <code>ctl_out.done</code> , <code>ctl_out.error</code> or <code>ctl_in.reset</code> .
<code>foo_any_written</code>	Writable and (array or sequence or string)	<code>bool_t</code>	Indication that any part of the value is being written. Persists until <code>ctl_out.done</code> , <code>ctl_out.error</code> or <code>ctl_in.reset</code> .
<code>foo_read</code>	Volatile or (readable and not writable)	<code>bool_t</code>	Indication that the property is being read. Persists until <code>ctl_out.done</code> , <code>ctl_out.error</code> or <code>reset</code> .

The *indication* signals are valid during the access operation (until `ctl_out.done` or `ctl_out.error` is asserted). The operation will be forced to complete with a timeout error if neither `done` nor `error` is asserted within the number of control clock cycles indicated in the `Timeout` attribute of the `ControlInterface` element. Unless the OWD declares that the `readError` or `writeError` attributes are true, *control software will not expect and not check that errors have occurred*. The `readSync` and `writeSync` OWD property attributes currently have no function for HDL workers.

Any `writable` property is registered in the worker's `shell` when written, even when the property is `volatile` and the worker is supplying a volatile value for reading in the `props_out` record. The signals possibly present in the `props_out` record are:

Table 11: HDL Worker Property Output Signals

Signals in <code>props_out</code>	Included when:	Type	Description
<code>foo</code>	Volatile or (readable and not writable)	*	The worker-supplied value of the property, with the type dependent on the property declaration.
<code>foo_length</code>	Volatile or (readable and not writable) and sequence type	<code>ulong_t</code>	The worker-supplied 32 bit unsigned length (number of elements in the sequence) when a sequence type.

3.7.6 Property Data Types

The `props_in` and `props_out` port signal records contain fields for property values with types that correspond to the property types defined in the OCS or OWD. All these types and associated conversion functions are defined in the `ocpi.types` package (in the `ocpi` HDL primitive library). This package is available in all workers. The worker author can decide to use the fully specified types (e.g. `ocpi.types.ulong_t`), or introduce the types into the worker architecture's namespace using:

```
library ocpi; use ocpi.types.all;
architecture rtl of xyz_worker is
```

For all property data types there is a:

- VHDL type name specified in the OCS with a `_t` suffix
- `from_<type>` conversion function from the type to `std_logic_vector`
- `to_<type>` conversion function from `std_logic_vector` to the type
- `to_<type>` conversion function from the related VHDL type (below) to the type
- `<type>_min` (for signed types) and a `<type>_max` constant for minimum and maximum values of the type
- `<type>_array_t` type for array or sequence property values, with a range of `(0 to length - 1)`
- `to_slv` conversion function from each `<type>_array_t` to `std_logic_vector`.
- `to_<type>_array` conversion function from `std_logic_vector` to `<type>_array_t`.

For example, for the `ushort` type, the `ocpi.types` package contains:

```
subtype ulong_t is unsigned (31 downto 0);
type ulong_array_t is array (natural range <>) of ulong_t;
constant ulong_max : ulong_t := x"ffff_ffff";
function to_ulong(c: natural) return ulong_t;
function to_ulong(c: std_logic_vector(31 downto 0)) return ulong_t;
function from_ulong(c: ulong_t) return std_logic_vector;
function to_slv(a: ulong_array_t) return std_logic_vector;
function to_ulong_array(a: std_logic_vector) return ulong_array_t;
```

The `string_t` type is a null-terminated array of `char_t` types. The `to_string` conversion function can convert from a VHDL `STRING` type to a `string_t`.

The types are summarized in the following table, with extra conversion functions specific to each type.

Table 12: VHDL Types for Properties

VHDL type	Based on	Width	Extra Conversion Functions
<code>uchar_t</code>	IEEE unsigned	8	<code>to_uchar(n : natural)</code>
<code>char_t</code>	IEEE signed	8	<code>to_char(i : integer)</code> <code>to_char(c : character)</code> <code>to_character(c : char_t)</code>
<code>ushort_t</code>	IEEE unsigned	16	<code>to_ushort(n : natural)</code>
<code>short_t</code>	IEEE signed	16	<code>to_short(i : integer)</code>
<code>ulong_t</code>	IEEE unsigned	32	<code>to_ulong(n : natural)</code>
<code>long_t</code>	IEEE signed	32	<code>to_long(i : integer)</code>
<code>ulonglong_t</code>	IEEE unsigned	64	<code>to_ulonglong(n : natural)</code>
<code>longlong_t</code>	IEEE signed	64	<code>to_longlong(i : integer)</code>
<code>float_t</code>	std_logic_vector	32	<code>to_float(r : real)</code> (not synthesizable)
<code>double_t</code>	std_logic_vector	64	<code>to_double(r : real)</code> (not synthesizable)
<code>string_t</code>	<code>char_t</code>	8	<code>to_string(s : string, length : natural)</code> <code>from_string(s : string_t) return string;</code>
<code>bool_t</code>	<code>std_logic</code>	1	<code>to_bool(b : boolean)</code> <code>to_bool(b : std_logic)</code> <code>its(b : bool_t) return boolean</code>

3.7.7 Raw Access to Properties

There is an alternative property access method when a worker must manage the storage and addressing of individual property values itself. This is called the **raw** property interface. There are two primary use cases for this method:

- Device workers using properties to access hardware registers outside the FPGA, e.g. via I2C, SPI.
- Application workers that need to arrange the storage of property values for more efficient storage of large values, e.g. in a block memory managed by the worker.

In both cases this avoids register duplication for property values, either off or on chip.

Raw access is enabled by setting the `rawProperties` attribute or the `firstRawProperty` attribute in the OWD. If `rawProperties` is true, then all properties are raw. If `firstRawProperty` names a property, then properties listed before that one are accessed using the mechanism described above, but the named property and later properties are accessed using the raw interface.

The input signals (in the `props_in` record) for the raw interface are:

Table 13: Raw Property Input Signals

Signal in <code>props_in</code>	Signal included when:	Signal Description
<code>raw.address</code> (31 downto 0)	Always	The byte offset from the first raw property, of the property being accessed.
<code>raw.byte_enable</code> (3 downto 0)	Some raw property is less than 32 bits.	The (4) byte enables for reading/writing bytes in the 32-bit data path of the control interface.
<code>raw.is_read</code>	Some raw property is readable/volatile	Access operation is reading a raw property, valid until <code>raw.done</code> or <code>raw.error</code> .
<code>raw.is_write</code>	Some raw property is writable/initial	Access operation is writing a raw property, valid until <code>raw.done</code> or <code>raw.error</code> .
<code>raw.data</code> (31 downto 0)	Some raw property is writable/initial	The data being written to the raw property, on the appropriate byte lanes for the offset.

The output signals (in the `props_out` record) for the raw property access interface are:

Table 14: Raw Property Output Signals

Signal in <code>props_out</code>	Signal included when:	Signal Description
<code>raw.data</code> (31 downto 0)	Some raw property is readable/volatile	The data value for the raw property being read, with values smaller than 32 bits (e.g. 8 or 16 bit values) aligned in the appropriate byte lanes. Valid and accepted when <code>raw.done</code> is asserted.
<code>raw.done</code>	Any raw properties	Indicates when the access cycle has completed successfully. Asserted for one cycle per access.
<code>raw.error</code>	Any raw properties	Indicates when the access cycle has completed unsuccessfully. Asserted for one cycle per access.

When raw properties are being accessed, the `props_out.raw.done` and `props_out.raw.error` signals indicate when the access is complete (and for reading, when the `props_out.raw.data` signal is valid). These are analogous to the `done` and `error` signals in `ctl_out`, although they do not have default values and *must be explicitly driven by the worker*.

If the raw interface is accessing registers or block memories in the worker, the `raw.done` signal may be tied asserted since all access happens in a single cycle. When the raw interface is used to access external registers (e.g. accessing an I2C or SPI bus), it would be asserted by the worker for one cycle when the access is complete.

3.7.8 Summary of HDL Worker Control Interface

- Interface inputs are in the `ctl_in` signal port record.
- Interface outputs are in the `ctl_out` signal port record.
- Clock is `ctl_in.clk`.
- Reset is `ctl_in.reset`, asserted high, *synchronously*, for at least 16 cycles.
- Do no work unless `ctl_in.is_operating` is asserted, or a control operation is in progress.
- Optionally use `ctl_out.done` and `ctl_out.error` when control operations or property accesses take more than one cycle.
- Optionally set `ctl_out.finished`, if the worker has some semantic of being *finished*.
- Property inputs (written values, and access indicators) are in `props_in`.
- Property outputs (volatile values) are in `props_out`.
- If raw properties are used, the interface is in the `raw` member of `props_in` and `props_out`.
- If using `ctl_out.done` or `ctl_out.error` (or their raw equivalents), and more than 16 control clock cycles are required to complete the operation, set the `Timeout` attribute in the `ControlInterface` element.

3.8 HDL Worker Data Interfaces for OCS Data Ports

Data interfaces convey messages with an associated **opcode**, which is an ordinal indicating the message type, among those defined in the protocol. When there is only one message type in the protocol, no opcode is used and no interface signals are present. Data interfaces convey boundaries between messages and thus messages have a well defined opcode, start, end, and length.

Data interfaces implement flow control, such that an output cannot be produced unless permission is granted. HDL workers explicitly accept data at input interfaces *when offered*, and respect permission to produce data at output interfaces.

The fundamental nature of data flow in and out of HDL workers is messages with opcodes, under flow control.

- Message oriented, with explicit boundaries between messages
- Messages are tagged with an opcode, indicating the type of message
- Input data is offered to and accepted by the worker
- Output data is provided only when permission is granted

A special use case for a data interface is when it has zero width. This is when all messages in the protocol have no arguments and thus are all zero length. This means that message opcodes are all that is conveyed. If there is only one message in the protocol, and it has no arguments, there is no data, no opcode, but still an indication of a message being conveyed. This is essentially an “event pipe” or “pulse” interface.

3.8.1 Message Payloads vs. Physical Data Width on Data Interfaces

Component ports defined in the OCS usually specify an OpenCPI protocol spec (OPS). This defines one or more messages that will be consumed or produced at the port. More details about protocols can be found in the CDG.

Each message payload has a serialized format as a sequence of bytes that, when used in software, are laid out in byte-addressed memory. For example, if the operation element in a protocol contains:

```
<argument name='a1' type='uchar' />
<argument name='a2' type='ushort' arraylength='2' />
<argument name='a3' type='ulonglong' />
```

And the values of this payload are:

```
a1: 1, a2: {0x2345,0x6789}, a3: 0xfedcba9876543210
```

Then the byte sequence (with proper alignment, and encoded little-endian), would be:

Sequence # ►	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Contents (hex)	01	x	45	23	89	67	x	x	10	32	54	76	98	ba	dc	fe
Argument	a1		a2[0]	a2[1]										a3		
Contents	1		0x2345	0x6789										0xfedcba9876543210		

This layout and these values is the same for all types of workers in all (little endian) environments and over all data paths. The **x** values are padding for alignment.

HDL worker data interfaces, associated with OCS ports, have a *physical* width, specified by the HDL-specific **DataWidth** attribute. It indicates the number of wires over which the messages will be conveyed. The width must be a multiple of the smallest data value in the protocol. In the example above this would be 8 bits. If **DataWidth** was 8, the sequence of content bytes shown above would be how the payload appears on that byte-wide data interface. If the **DataWidth** was 16, the message would appear as:

Sequence # ►	0	1	2	3	4	5	6	7
15 downto 8	x	23	67	x	32	76	ba	fe
7 downto 0	01	45	89	x	10	54	98	dc

If the **DataWidth** was 32, the message would appear as:

Sequence # ►	0	1	2	3
31 downto 24	23	x	76	fe
23 downto 16	45	x	54	dc
15 downto 8	x	67	32	ba
7 downto 0	01	89	10	98

And if the **DataWidth** was 64, the message would appear as:

Sequence # ►	0	1
63 downto 56	x	fe
55 downto 48	x	dc
47 downto 40	67	ba
39 downto 32	89	98
31 downto 24	23	76
23 downto 16	45	54
15 downto 8	x	32
7 downto 0	01	10

The byte sequence remains the same regardless of **DataWidth**.

3.8.2 Byte Enables on Data Interfaces

Byte enables on data interfaces are only present when needed, and their presence is determined by a combination of the protocol and the **DataWidth** of the interface.

Two values are inferred from the protocol:

- **DataValueWidth**: the smallest data value in the protocol.
- **DataValueGranularity**: the least common multiple of data values among all messages in the protocol; all message lengths are a multiple of this number of data values.

The physical data width of the interface, **DataWidth**, must be a multiple of **DataValueWidth**. When **DataWidth** is greater than **DataValueWidth** * **DataValueGranularity**, byte enables are in the interface, since data words (of **DataWidth**) at the end of a message may be partially valid. E.g. if the **DataWidth** is 32, and **DataValueWidth** is 8 and **DataValueGranularity** is 1, the messages may have any number of bytes and thus the last 32 bit word of a message may have 1, 2, 3 or 4 valid bytes. In this context, a byte is a data value, and bytes *may* not be 8 bits. Here are some examples:

- Message is a sequence of short (16 bit) values, **DataWidth** is 16:
DataValueWidth = 16
DataValueGranularity = 1
No byte enables required.
- Message is a sequence of short (16 bit) values, **DataWidth** is 32:
DataValueWidth = 16
DataValueGranularity = 1
Byte enables (2) are required since sequences might be an odd number of shorts.

- Message is a sequence of *pairs* of short (16 bit) values, **DataWidth** is 32:
DataValueWidth = 16
DataValueGranularity = 2
Byte enables not required since sequences are always a multiple of 2 shorts.

The exact naming and use of byte enable signals in data interfaces is described below.

3.8.3 Streaming Data Interfaces to/from the HDL Worker

Worker data ports (as specified in the OCS) can be implemented in two different styles: stream or message. Stream interfaces are FIFO-like with extra qualifying bits along with the data for message boundaries and byte enables. Message interfaces are based on addressable message buffers, and are described in the following section [Message Data Interfaces](#).

The stream interface is the default style for data ports. If all attributes of the interface have default values, no indication of this style is required in the OWD.

Streaming data interfaces must implement message boundaries and flow control, described in detail in the [Metadata and Message Boundaries](#) and [Flow Control](#) sections below. Message boundaries mean that input and output data is delimited such that the stream of data has markers indicate when messages start and end. Flow control means:

- Data must not be produced (given) unless permission is granted (output it *ready*).
- Data must not be consumed (taken) unless permission is granted (input is *ready*).

When data port attributes are required, this style is indicated by including a **<StreamInterface>** XML element in the OWD.

An example of the per-data-port XML element is:

```
<StreamInterface name="sensor"
                 dataWidth="64"
                 preciseBurst="true"/>
```

The example shows the **sensor** data interface port declared in the OCS is being further configured for a non-default **dataWidth** and support for **preciseBurst**.

3.8.3.1 Implementation-Specific Attributes of Streaming Data Interfaces

Recall that the **dataWidth** attribute at the top level of the OWD specifies the default physical width of all data interfaces. If most interfaces have the same width, it may be most convenient to specify this width at the top level **HdlWorker** element.

DataWidth specified in the **StreamInterface** element applies only to that port.

Other than **DataWidth**, there are two options that can be specified for a stream data interface: **precise bursts** and **aborts**.

Precise bursts are messages where the length of the message is available at the *start* of the message. Since the message is being streamed through a fixed width interface, the default mode is to terminate the message with a marker at the *end*. Producing precise bursts has value in that the receiver can take advantage of the early knowledge

of the message length, but precise bursts may have a cost in latency or storage when data must be buffered to determine the length of the message.

Declaring precise bursts at an input interface means that the worker *requires* precise bursts: it must know the length of a message at the start of the message. Declaring precise bursts at an output interface means that the worker will produce precise bursts. If a worker requiring precise bursts on input is connected to a worker that does not produce precise bursts on output, an adapter worker must be inserted which converts imprecise bursts to precise bursts.

Precise bursts are an optimization that allow use of message lengths that are available early, such as when a message is arriving from a network with a header that specifies its length.

Aborts are used when a worker producing a message finds that it must interrupt the message it is producing in such a way that the part that was already sent may be invalid. Declaring aborts at an input interface means that the worker can *tolerate* aborted messages and act accordingly. Declaring aborts at an output interface means that the worker may indeed abort messages after starting but before ending a message.

Aborts are a latency optimization such that a consumer can start processing a message before knowing if it will complete without aborts. This is used in network processing to reduce average latencies. An example is when a CRC check at the end of a network message will determine whether the message is valid or not. Waiting until the CRC is checked before processing a message adds latency in the common case where the CRC is good.

If a worker that produces aborts is connected to a worker that can not tolerate aborts, a buffering adapter worker must be inserted which buffers the message and discards it if it is aborted.

The XML attributes of the `StreamInterface` element are in the following table:

Table 15: XML Attributes of StreamInterface Elements

Attribute in StreamInterface	Attribute data type	Attribute Description
<code>name</code>	string	The name of the port in the OCS. Required.
<code>dataWidth</code>	unsigned	The width of the data path for this interface. The default is the smallest element in the message protocol indicated in the OCS, unless overridden by a default <code>datawidth</code> attribute at the top level of this OWD (<code>HdlWorker</code>)
<code>preciseBurst</code>	boolean	Whether this interface requires (on input) or will produce (on output) precise bursts. The default is false. If a worker can easily support precise bursts on output, it should.
<code>abortable</code>	boolean	On input, that aborted messages are tolerated. On output, that the worker may abort messages.
<code>pattern</code>	string	Signal naming pattern. This applies to external signals, and not the inner worker signals. The default is "%s_ ", which is simply the port name as prefix with an underscore.
<code>clock</code>	string	The name of the clock used with this port. The name is either the name of another port with a "myclock" attribute of true, or the name of a clock declared for the worker as a whole, that is not the control clock. (Not supported in current framework)
<code>myclock</code>	boolean	An indication that this port operates in its own clock domain, and will have its own <code>clk</code> signal as input. (Not supported in current framework)

Although a number of attributes of data interfaces are inferred/derived from the message protocol specified in the OCS/OPS, the OCS may not have a message protocol or certain attributes may need to be overridden. Below are the protocol-related attributes that can be specified in a `StreamInterface` element.

The indicated defaults apply only when there is no message protocol in the OCS.

Table 16: Stream Interface Attributes for Overriding Protocol Defaults

Attribute in StreamInterface	Attribute data type	Attribute Description
		<i>The indicated defaults only apply if there is no message protocol in the OCS. Otherwise a value is inferred from the protocol and no default applies.</i>
<code>numberOfOpCodes</code>	unsigned	The number of distinct message types (opcodes) for the interface. The default is 1, the maximum is 256.
<code>maxMessageValues</code>	unsigned	Largest number of data values in a message. Default is 64K.
<code>dataValueWidth</code>	unsigned	The size of the atomic units of data on this interface (size of a byte). The default is the value of the <code>datawidth</code> attribute.
<code>dataValueGranularity</code>	unsigned	The smallest multiple of number of data values that will actually appear in any message.
<code>zeroLengthMessages</code>	boolean	Whether the interface will support zero length messages. Default is false.

Specifying these attributes in the `StreamInterface` element determines the signals present in the input and output records for the interface.

3.8.3.2 Metadata and Message Boundaries Used for Stream Interfaces

The information flowing out of or in to stream interfaces are variable length messages conveyed using fixed width **words**, which are a combination of data and metadata. The metadata associated with every word presented at the interface includes:

- **SOM**: start of message: indicates that this word is the first in a message, *whether or not there is data present*.
- **EOM**: end of message: indicates that this word is the last in a message, *whether or not there is data present*.
- **Valid**: indicates whether the data part of this word has valid data
- **Abort**: (optional) whether this word is indicating the end of an aborted message
- **Byte_enable**: (optional) indicates, *when valid is true*, which bytes in the data word are valid. This signal is all ones on all but the last valid word of a message.

These metadata signals, as well as the `data` signal, are *all* registered (at least x1) on input, outside the worker's code, enabling simple workers to be written in a combinatorial style.

On input, the optional **Abort** indicator, when present and asserted, forces the **EOM** indicator on, and the **Valid** indicator off. When **Abort** is asserted on output, **EOM** and **Valid** are ignored and assumed true and false, respectively. When there is no **Abort**,

the three metadata bits, **SOM**, **EOM**, and **Valid**, can be in any combination, with the following meanings:

Table 17: Metadata in Stream Interfaces

SOM	Valid	EOM	Signal Description
1	0	0	The start of a message, without any associated data. Sometimes called an “early” start of message.
1	1	0	The start of a message, coincident with the first word of data.
1	0	1	A zero length message, with no data, in a single word.
1	1	1	A single word message.
0	0	0	Reserved
0	0	1	A trailing end of message, with no data. Sometimes called a “delayed” or “late” end of message.
0	1	0	A data value in the middle of a message
0	1	1	A data value, coincident with the end of the message

On input, the worker can assume that the sequence of metadata values will be correct, meaning all messages will have a SOM and an EOM. The valid sequences for conveying messages *will* follow these rules on input and *must* follow them on output:

- The first word of a message must have **100**, **110**, **101**, or **111** (i.e. have the **SOM** bit set).
- After a message is started but not simultaneously ended by **101** or **111**, any of **001**, **010**, or **011** may occur (i.e. be data-without-**EOM**, data-with-**EOM** or **EOM**-without data).
- After a word with **EOM** set, the next word must have **SOM** set.

Output ports *must* limit the size of messages by using **SOM** and **EOM**. Usually this involves having output messages be sized and delimited to track the size of messages arriving at input ports. Sometimes a worker must specifically determine and implement output message sizes based on some other criteria. An example would be a worker that produced a fixed size message regardless of the size of input messages, essentially accumulating data from input messages into fixed size output messages.

The data and metadata flowing through an interface is qualified by an input **ready** signal at both input and output interfaces. When this signal is not asserted (from outside the worker), the data and metadata signals are *invalid* at input and *ignored* at output.

3.8.3.3 Flow Control (a.k.a. Back Pressure) On Streaming Interfaces.

Both input and output stream interfaces have a **ready** signal, that is *always input to the worker*, indicating that data can be consumed or produced.

The rules for *input* interfaces are:

- **ready** indicates that the metadata and perhaps the data signals are valid
- if **ready** is not asserted, none of the metadata signals are valid or meaningful
- the worker **takes** input data when **ready** is asserted by asserting the **take** signal
- it is invalid to assert the **take** signal if the **ready** input signal is not asserted

The rules for *output* interfaces are:

- **ready** indicates that metadata and perhaps data can be produced
- if the **ready** signal is not asserted, none of the metadata or data output signals are considered valid.
- the worker **gives** data when **ready** is asserted by asserting the **give** signal
- It is invalid to assert the **give** signal if the **ready** input signal is not asserted

Data flows according to FIFO semantics. Input data is presented as **ready** as if there is an input FIFO outside the worker that is **not empty**. The worker consumes this data by **taking** it, which is as if it is **dequeueing** data from this **not-empty** FIFO. Similarly, output data can be produced when output is indicated to be **ready** as if there is an output FIFO outside the worker that is **not full**. The worker produces this data by **giving** it, which is as if it is **enqueueing** data to this **not-full** FIFO.

These signals (**ready**, **take**, **give**) control the flow of data and metadata words through the interface. Here is a table of how this signal terminology compares to some other common interfaces with FIFO semantics: the classic FIFO interface, the AXI streaming interface, and the “native” Xilinx FIFO interface.

Meaning	OpenCPI	Classic FIFO	AXI	Xilinx FIFO
Data is available to consume	ready	not_empty	valid	!empty
Consume data	take	dequeue	ready	rd_en
Data can be produced	ready	not_full	ready	!full
Produce data	give	enqueue	valid	wr_en

In AXI interfaces, either signal (**valid** or **ready**) may be asserted early. The handshake (**ready**) can in fact be asserted early even when **valid** is not yet asserted. With OpenCPI it is invalid to assert **take** or **give** without **ready** being asserted. In Xilinx FIFO, **rd_en** and **wr_en** are ignored if the fifo is **empty** (input) or **full** (output).

In all cases (all these interfaces) data moves from producer to consumer when both sides assert their signals in the same cycle (at the same clock rising edge).

Since worker ports all have FIFO semantics, workers must be written to accommodate “back pressure”. I.e. the **ready** signal on output interfaces may not always be asserted, so the output data must be held/stopped until it can be **given**.

Example timing diagrams for this interface follow the signal descriptions below.

3.8.3.4 Signal Summary for Streaming Interface

The signals for stream interfaces (input or output) are in the following table.

Table 18: Stream Interface Input and Output Signals

Signal	When Included	Direction	Signal Description
clk	When myClock or Clock	Input	The clock for this interface, only when the interface uses a specific clock. Otherwise the <code>ctl_in.clk</code> signal should be used.
ready	Always	Input	On an input port: can consume, on an output port: can produce. On input metadata, bits are valid. Worker can assert take (input) or give (output)
data	If datawidth > 0	Data direction	The input or output data, when valid is true.
som	Always	Data direction	The start-of-message indication. For output ports, if previous give indicated eom , then this som indication is assumed true on the next give.
eom	Always	Data direction	The end-of-message indication is present
valid	If datawidth > 0	Data direction	The data signals hold message data.
abort	If Abortable	Data direction	The message is being aborted
byte_enable	**	Data direction	Which data bytes are valid; usable when valid is true. **Included only when the WithValueWidth * DataValueGranularity < datawidth .
opcode	numberOfOpCodes > 1	Data Direction	The opcode for the current message. Valid from start of message to end of message on input. Must be valid with som on output.
give	Port is output	Output	Indicates word is given to output port, allowed when ready is true. Like an enqueue signal to a FIFO. When asserted, at least one of som , eom , valid , or abort must be asserted.
take	Port is input	Output	Indicates that word is being taken by worker, only permissible when ready is true. Worker is taking the data this cycle. Analogous to dequeuing a FIFO.

The `data` and `byte_enable` signals are `std_logic_vector`. The opcode signal is also `std_logic_vector` when there is no protocol in the OCS, otherwise it is an enumeration of type `<protocol>_OpCode_t`, with each operation having an enumeration constant `<protocol>_<op>_op_e`. These opcode types are in the `work.<worker>_worker_defs` package. The other signals are all `bool_t`, from the `ocpi.types` package.

3.8.3.5 Timing Diagrams

The following diagram shows an input port where both the infrastructure (worker shell) and the worker respond one cycle after they see new input, resulting in a throughput of 3 clock cycles per data word. Both SOM and EOM are coincident with data.




Figure 4: 3 word input message with delays on both sides

The following diagram shows an input port where only the worker responds one cycle after it sees new input, resulting in a throughput of 2 clock cycles per data word.




Figure 5: 3 word input message width worker adding delay

The following diagram shows an input port where both sides respond in the same cycle, with no delays, resulting in a throughput of 1 clock cycles per data word.




Figure 6: 2 3 word messages, with no delays on either side.

3.8.3.6 Source Code Examples

Here is a set of examples, showing the use of these signals. The first is a complete worker that is purely combinatorial, and simply adds a constant to every data value from input to output. No VHDL process or clocking or even reset is used since `ctl_in.is_operating` reflects the reset condition, and the computation takes place in a single clock cycle. No opcode or byte_enable is used since the protocol has a single operation.

```
architecture rtl of example_worker is
  signal doit : bool_t;
begin
  doit      <= ctl_in.is_operating and in_in.ready and out_in.ready;
  in_out.take <= doit;
  out_out.give <= doit;
  out_out.data <= std_logic_vector(unsigned(in_in.data) + 3);
  out_out.som <= in_in.som;
  out_out.eom <= in_in.eom;
  out_out.valid <= in_in.valid;
end rtl;
```

The next example shows a worker that inserts a special single-word all-ones message, with operation `xy`, every 8 messages it passes through. The OCS protocol is `myprot`.

```

architecture rtl of example_worker is
    signal count      : unsigned(3 downto 0);
    signal ready      : bool_t := ctl_in.is_operating and out_in.ready;
    signal inserting  : bool_t := count = 8;
begin
    process (ctl_in.clk) is
    begin
        if rising_edge(ctl_in.clk) then
            if its(ctl_in.reset) then
                count <= (others => '0');
            elsif its(ready) then
                if its(inserting) then
                    count <= (others => '0');
                elsif in_in.ready and in_in.eom then
                    count <= count + 1;
                end if;
            end if;
        end if;
        end process;
        in_out.take      <= ready and in_in.ready and not inserting;
        out_out.give     <= ready and (count = 7 or in_in.ready);
        out_out.data     <= (others => '1') when inserting else in_in.data;
        out_out.som      <= inserting or in_in.som;
        out_out.eom      <= inserting or in_in.eom;
        out_out.valid    <= inserting or in_in.valid;
        out_out.opcode   <= myprot_xy_op_e when inserting else in_in.opcode;
    end rtl;

```

3.8.4 Message Data Interfaces

[Unsupported in 2017Q2]

This interface is used when an OWD specifies a MessageInterface element associated with a data port in the OCS. It provides an alternative mechanism to consume or produce data that uses an addressable message buffer interface. This enables the worker to produce or consume message data out of order, or to only access parts of a message. The signals provided with this interface allow the worker to address specific locations in the *current buffer*, and then signal that it is done with the current buffer.

3.8.5 HDL Worker Data Interface Summary

- Only move data when `ctl_in.is_operating` is true.
- Input data is valid only when explicitly indicated at the input interface.
- Output data can only be moved when flow control allows it at the output interface.
- Synchronize control signals if data interface clocks are different than control clock.
- Output message boundaries must be supplied and must respect maximum message sizes.
- Output message boundaries are usually derived from input message boundaries.

- Remember to (in most cases) convey zero length messages from input to output.

3.9 Time Service Interface

This interface provides “time of day” information to the worker, to the precision requested in the OWD via attributes to the **TimeInterface** element. Time of day values are supplied to the worker in the clock domain of this interface, which defaults, like all interfaces, to the control clock.

The signals for the time service interface are in the **time_in** signal port record and are described below. If the default port name is overridden, the signal port record could be **<port-name>_in**.

Table 19: Time Service Signals

Signal in time_in	Width	Signal Description
seconds	SecondsWidth attribute of TimeService element	The entire seconds part of the time-of-day, in GPS time (no leap seconds). If the width is 32 it is absolute time. If width less than 32, it is just a relative time truncated preserving the LSB, to that value, and wraps. The LSB is always 1 second; VHDL type is IEEE numeric unsigned. Width may be zero, in which case this signal is not present.
fraction	FractionWidth attribute of TimeService element	The binary fraction of a second, with the radix point to the left of the MSB. If width is 32 bits, the LSB represents 2^{-32} seconds, or ~233 ps. If width is less than 32, the MSB are preserved, such that the MSB is always $\frac{1}{2}$ second. Width may be zero, in which case this signal is not present.
valid	1 (bool_t)	Indicates when the time of day is valid. Present only when the AllowUnavailable attribute is true.

3.10 Memory Service Interfaces

This interface provides access to memory. [Not supported in 2017Q1]

4 Building HDL Assets

Building workers is similar across different authoring models and languages: typing the `make` command in a worker, library, or project directory builds workers for a specified list of targets. If the worker depends on lower level “primitive” libraries, those libraries are referenced in the worker's `Makefile`, using the `Libraries` or `Hdllibraries` variable. To make a declaration of such primitive libraries common to all workers in a library, or in a project, the `Hdllibraries` variable can be set in the library's `Library.mk` file or in the project's `Project.mk` file. This variable applies only to HDL workers. Lower level libraries must be built before building a worker which references them.

Similar to other authoring models (e.g. the RCC model), the build targets are specified by setting variables specifying targets for that model. For HDL, these target variables are: `Hdltarget`, `Hdltargets`, `Hdlplatform`, and `Hdlplatforms`. The singular form specifies one target, and the plural forms specify multiple targets separates by spaces. HDL build targets discussed in detail in the next section.

The default target for RCC workers is the development system itself (e.g. `linux-c7-x86_64`). Unlike RCC workers, HDL assets have no inherent default target. However, a default value for the various HDL target variables can be set in the worker `Makefile`, the library's `Library.mk` file, or in the project's `Project.mk` file.

For software workers, this is usually the end of the build process: deployable artifacts for these workers are created and ready for export and/or use in applications.

For HDL workers, it is different. FPGAs are generally not subject to dynamic, partial loading: the whole FPGA must be reloaded with a full “configuration bitstream”. [OpenCPI does not support *partial reconfiguration* of FPGAs as of this document version]. As with any authoring model, primitives are built first. Then HDL workers are built and, for targets that are real FPGAs rather than simulators, synthesized. Finally, there are two additional steps in the build process in order to create the final, dynamically loadable configuration bitstream:

- Composing workers into an **HDL assembly**.
- Finalizing the bitstream as an **HDL container**.

This final step creates the deployable artifact usable for export and/or use in applications. These steps are defined in the sections below.

4.1 HDL Build Targets

Build targets specify the target device, family of devices, or platform for which the asset should be built (compiled, synthesized, place-and-routed, etc.) When building any level of modules for FPGAs, the build targets are specified via the `Hdltarget(s)` or `Hdlplatform(s)` variables. The *targets* are chips or chip families, whereas the *platforms* are actual FPGAs on specific boards. HDL primitives, workers, and assemblies, are built for `Hdltargets`, and HDL containers (final bitstreams) are built for `Hdlplatforms`. These build targets are defined in a hierarchy with these levels:

Top level, vendor level: this level specifies vendors (Xilinx, Altera), as well as vendor-independent simulators (Modelsim). This enables HDL assets to be built for all Xilinx and Altera targets or built for Modelsim. This implies building for all lower level targets under these top-level labels. The value `all` specifies all top level supported targets.

Family level: this level specifies the family of parts under the vendor level. Different part families typically have different on-chip architectures, and may drive tools differently. Building for a family target means generating libraries or cores that are suitable to any member (part) in the family. Examples would be “virtex6” or “zynq” or “stratix4”. Simulation targets at the top level don’t have families (yet) so these top two levels are the same for simulation.

Part level: specifies the exact part the design is targeted at, e.g. `xcv5lx50t`. This does not include package information but may include speed grades.

The following two `make` variables can further filter the targets that are built anywhere that HDL building takes place.

ExcludeTargets/ExcludePlatforms: this variable specifies targets to be excluded, usually because they are known not to be buildable for one reason or the other (a tool error, or other incompatibility).

OnlyTargets/OnlyPlatforms: this variable specifies targets to be exclusively included, because it is known that only a limited set of targets should be built (e.g. a Xilinx coregen core specific to a particular family or part).

The `HdLPlatforms` variable specifies HDL platforms (like Xilinx `m1605` and `zed`), which imply the appropriate family and part. I.e., if you specify to build for a platform, it will build primitives and workers for the appropriate part family. Except for the final bitstream build, the `HdLTarget(s)` are implied by the `HdLPlatform(s)`.

If no HDL target variables are set, the `OCPI_HDL_PLATFORM` environment variable can be set to an HDL platform, and that will be used for all HDL builds.

In some synthesis cases, tools that target a *part family* in fact target the *smallest* part in the family and try to limit use of some on-chip resources (e.g. DSP blocks) to the amount that exists on the smallest part. While this usually correctly generates a resulting file that can be used on any part in the family, it is not always desirable when the target *platform* in fact has a *larger* part. To force the use of a specific part for an `HdLTarget`, in the HDL worker `Makefile`, insert a line setting the `HdLExactPart` variable, e.g.:

```
HdLExactPart=\$ (and $(filter virtex6,$(HdLTarget)))
```

If a worker is being built for only one target in any case, then this would work:

```
HdLExactPart=xc6vlx240t
```

On the command line this could be:

```
% make HdLTarget=virtex6 HdLExactPart=xc6vlx195t
```

4.2 The HDL Build Hierarchy

OpenCPI FPGA bitstreams (the files that configure entire FPGAs) are built in several layers. The same layers apply to building executables for simulation. The following diagram shows the build flow (bottom to top) and hierarchy.




Figure 7: OpenCPI HDL Build Flow Layers

At the bottom layer (built first, used by all other layers), are **primitives**. These are low level, “leaf” libraries and cores used by higher levels. **Primitive libraries** are libraries of modules built from HDL source code that are available to be used higher up the hierarchy; using a primitive library in a higher level module does not imply all the modules in the library are brought into the design, but only pulled in as needed by references in the higher levels of the design.

Primitive cores on the other hand are single modules built from source or generated from tools such as Xilinx Coregen, which are also used in higher levels of design. They are explicitly included in workers. Primitives may in fact depend on each other: a core may depend on primitive libraries, and primitive libraries may depend on other primitive libraries. Circular dependencies are not supported.

There are primitive libraries specific to vendors and families that can be used for implementing primitives using vendor-specific elements. More detail on creating such primitive libraries are in the **OpenCPI HDL Platform Development** document.

Above the **primitives** layer is the HDL **worker** layer, with workers of several types. All types of workers can use primitive libraries or cores as required. Application workers are generally portable and hardware independent. Device workers are workers that connect to the I/O pins of external hardware, and in some cases can attach to vendor-specific on-chip structures (e.g. ICAP on Xilinx). Adapter workers are used when two connected workers are not connectable in some way due to different interface choices in the OWD (e.g. width, stream-vs-message, clock domains). Adapter workers are normally inserted automatically as needed.

A platform worker is the special type of device worker that performs necessary platform-wide functions for the platform.

At the next layer, the HDL **assembly** is automatically generated HDL source code that uses application workers and adapter workers. The HDL assembly itself is described in metadata (XML) as an assembly of connected application workers. It typically represents a subset of an overall heterogeneous OpenCPI application: a subset that will be executed on a single FPGA.

The **platform configuration** is automatically generated HDL source code that uses platform workers along with some device workers. It represents a platform configured with built-in support for some attached devices, and may include various constraints and physical design. For those familiar with Linux kernels, a platform configuration is analogous to a built/configured kernel with some device drivers built-in.

At the top layer, the **container** adapts the application assembly to a platform configuration and provisions any additional required device workers. It connects and adapts the “external I/O ports” of the HDL assembly to the available I/O paths and devices in the platform. When the deployment of the HDL assembly requires device workers that are not in the platform configuration, they are instanced in the container itself. The Linux kernel analogy is that these extra device workers are analogous to the dynamically loaded device drivers used to run the application.

Device workers can either be built into the platform configuration or instanced in the container.

The final design for the entire FPGA is the container logic. This hierarchy (except primitives) is shown in the following diagram.




Figure 8: HDL Full FPGA Hierarchy

When tools support it, each layer in the build is actually synthesized or precompiled or elaborated as the tools allow (i.e. “prebuilt”).

- Each worker in a component library is prebuilt (possibly using primitive libraries and cores)
- Assemblies are prebuilt from generated VHDL or Verilog code using the required worker cores
- Platform configurations are prebuilt from platform workers and device worker cores
- Container top levels are built from platform configurations and HDL assemblies, with any additional device workers, service modules and adapters as required
- Full bitstreams (or simulation executables) are built from the container modules

This layered prebuilding allows the results to be reused at the next level without recompiling or resynthesizing, all in a vendor independent fashion. E.g. an HDL assembly prebuilt for a Xilinx virtex6 part can be reused to target different virtex6-based

platforms. The exact definition of prebuilding varies with different tool chains, and the level of synthesis optimization that happens at each step also varies by tool, and some of this level of hardening at each level is controllable for some tools.

At one extreme, prebuilding simply means remembering which source files must be provided to the next level (for tools that have no precompilation of any kind). At the other extreme are tools that can incrementally synthesize to relocatable physically mapped blocks on a family of FPGA parts.

Simulators are considered HDL platforms that act as test benches for assemblies. This is described in more detail below in the simulation section.

4.3 HDL Directory Structure

In OpenCPI projects, HDL workers are in component libraries with other non-HDL workers. Component libraries are thus heterogeneous, where different workers, possibly using different authoring models, may implement a common spec. A project may have a single component library called 'components', where worker directories are located. Alternatively, a project may have multiple uniquely named libraries under the `components` directory, each of which contains workers with some common theme. When projects have a single component library, it is in the `components` directory of the project. There can also be multiple separately named component libraries *under* the `components` directory. Additionally, the top-level `hdl` directory contains the following directories:

primitives: This directory contains subdirectories for each primitive library or core.

assemblies: This directory contains subdirectories for each HDL assembly of application workers, and is where containers deploying these assemblies on platforms are built into bitstreams and simulation executables

devices: This directory is a component library containing HDL device worker for devices that are potentially usable on different platforms. HDL device emulators and software proxies for some of the devices may also be in this component library.

platforms: This directory contains subdirectories for each platform implemented in the project. Platforms are either a specific FPGA chip/part on a circuit board with attached devices, or simulators. This is where platform-specific worker code exists, and where platform configurations are specified and built. There may also be a subdirectory under the platform's directory, called `devices`, containing a library of HDL device workers, proxies and emulators specific to that platform.

cards: This directory contains HDL device workers (and their proxies and emulators) that are specific to cards, rather than those generally useful on different platforms and cards. It also contains specification files for cards.

Development for HDL devices, platforms, and cards is described in the [*OpenCPI Platform Development Guide*](#).

Application workers for all authoring models are found in component libraries, which are either standalone outside of projects, in projects, or part of the [*OpenCPI CDK*](#).

This directory hierarchy is shown in the following diagram. All directories are optional and are created as needed by the `ocpidev` tool described in the CDG.

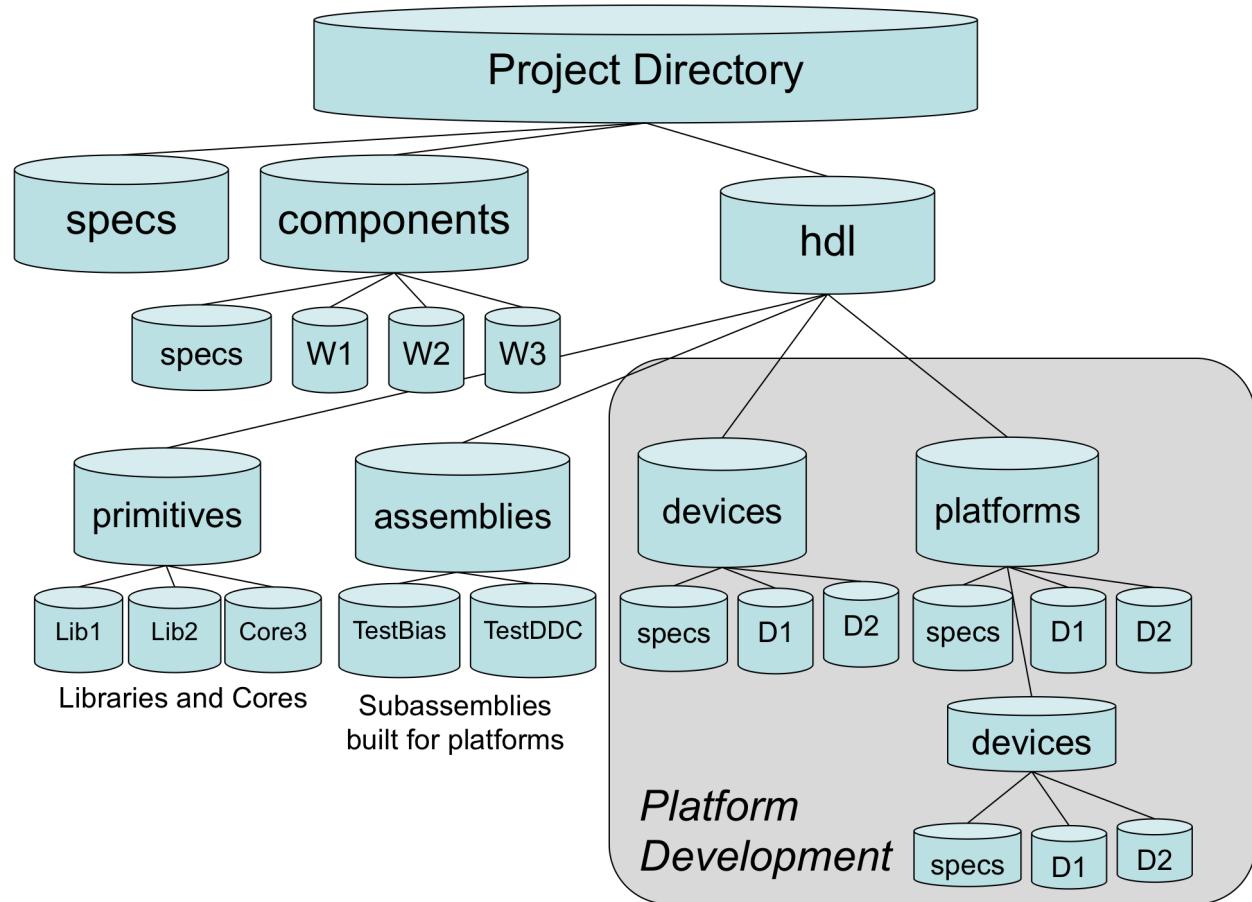


Figure 9: OpenCPI HDL Directory Structure

4.4 HDL Search Paths when Building

When building at any level of the HDL build hierarchy, the asset being built can depend on other assets, usually at lower levels of the hierarchy. HDL primitive libraries and cores can depend on other HDL primitive libraries. Workers can depend on primitives, specs, and protocols. HDL assemblies always depend on application workers in libraries. The assets being depended on may be in the same library or project, or in a different project, or in the CDK itself.

In each case of a dependency, the underlying asset is found using search rules. The built-in search rules automatically find assets that are in the component library, when building workers in that library. They also automatically find assets that are in the current project, when building HDL primitives, workers, and assemblies in projects. Special action by the developer is only required when there are dependencies on assets outside the current library or project.

If assets in one project (A) depends on assets in another (B), that should be stated in the `ProjectDependencies` variable in the project A's `Project.mk` file. If a special version of that second project (B) was needed temporarily, an environment setting would cause the special version to be searched first, shadowing the assets in the project (B) explicitly stated in `ProjectDependencies` of A. By “shadowing” we mean that the environment setting causes the search to look elsewhere *before* looking at the projects specified in the `ProjectDependencies` variable.

Environment variables are normally used only to temporarily replace/shadow assets in the default search path. Dependencies are normally stated at the point of dependency and stored in the appropriate file that is part of that asset (e.g. `Project.mk` for project-level dependencies). Environment variables essentially override these stated dependencies.

For each of the search rules defined below, the following principles are applied:

- Search “closer” first, then “farther away” (e.g. in library, then in project, then other projects, then CDK).
- Search within the project before searching outside the project.
- Search using an asset-specific path specified in the environment, before searching other projects.
- Search using the project path in the environment (`OCPI_PROJECT_PATH`) before using the explicit dependencies in the project's `Project.mk` file.
- Search the CDK last.
- Search the directories specified in any path environment variables in the order they appear in the colon separated list.

4.4.1 Searching for HDL Primitives

When a worker depends on primitive libraries, it specifies this by declaring the library name in a list in the `Libraries` variable in its `Makefile`. Similarly, when an HDL

worker depends on primitive cores, it specifies this by putting the core name in a list in the `Cores` variable. When the worker is being built, these lists are used in the following way.

If the library or core name has slashes in it, it is treated as a pathname (absolute or relative) to the specific directory of the primitive library or core. Examples are:

```
Libraries=../../myprims /home/colleague/hisprims/funprims  
Cores=../../..../ourcores/fft
```

If a name in the primitive library or core list does *not* have a slash, it is found by searching in the following places, in the following order:

- The HDL primitive libraries or cores in the current project (if the worker is in a project).
- The HDL primitive libraries or cores in the directories listed in the environment variable `OCPI_HDL_PRIMITIVE_PATH`, which are colon separated and searched in order.
- The HDL primitive libraries/cores in the projects listed in the environment variable `OCPI_PROJECT_PATH`, which are colon separated and searched in order.
- The HDL primitive libraries or cores in the projects listed in the `ProjectDependencies` variable in the project's `Project.mk` file.
- The HDL primitive libraries or cores in the *OpenCPI CDK*, identified in the environment variable `OCPI_CDK_DIR`.

When it is convenient to put a list of primitive libraries in the library `Makefile` (making them available to all workers in the library), the name of the variable is specific to the authoring model. Thus to make a list of libraries available to all the HDL workers in a component library, you would put the following line in the *library's Makefile*:

```
HdlLibraries=/home/george/primlibs/gprims
```

4.4.2 Searching for XML files (OCS, OPS) when Building Workers

As described in the CDG, all workers have an OWD, and all OWDs depend on a component spec, normally found in a separate OCS XML file. Furthermore, OCS files frequently depend on a separate OPS (protocol spec) file. It is also possible that an OWD could include an XML file to incorporate a list of properties defined elsewhere.

When looking for these XML files, when their name has no slashes, it is found by looking in the following places, in order:

- The worker's directory
- The `gen` subdirectory of the worker's directory
- Directories specified in the space-separated list in the make variable `XmlIncludeDirs`
- The component library's export directory (`lib/hdl`) for referencing other HDL workers (e.g. for `slave` and `emulate` attributes)

- The component library's **specs** subdirectory
- The export directories (**lib/hd1**) of all component libraries in the **ComponentLibraries** list (see below) for referenced HDL workers (e.g. for **slave** and **emulate** attributes)
- The directories listed in the environment variable **OCPI_XML_INCLUDE_PATH**, which are colon separated and searched in order
- The **specs** subdirectories of all projects in the environment variable **OCPI_PROJECT_PATH**, which are colon-separated and searched in order
- The **specs** subdirectories of the projects listed in the **ProjectDependencies** variable in the project's **Project.mk**
- The **specs** directory of the **OpenCPI** CDK, located using the environment variable **OCPI_CDK_DIR**

4.4.3 Searching for Workers in Component Libraries

The **Makefile** variable **ComponentLibraries** specifies a list of places to look when searching for workers. The most common use case for **ComponentLibraries** is for creating HDL assemblies, where the workers specified in the assembly must be found by searching for them in component libraries. Other uses include device workers, platform workers and platform configurations, as described in [**HDL Platform Development**](#).

While it might seem counter-intuitive for a worker *inside* a component library to depend on other component libraries, there are three cases where this occurs:

- A worker's OWD depends on specs (OCS and/or OPS) in another library.
- A worker is a proxy for a worker defined in another library (HDL workers cannot act as proxies).
- A worker is a device emulator for a device worker defined in another library.

In all these cases, a worker or a component library might define the **ComponentLibraries** variable to specify this dependence.

To find workers, the search first looks in the library that the worker is already a part of. After this, the **ComponentLibraries** variable is used, which holds a list of component libraries to search.

If the component library name in **ComponentLibraries** has slashes in it, it is treated as a path name (absolute or relative), to the specific directory of the component library. If a name in the list does *not* have a slash, the component library is found by looking in the following places, in order:

- The component libraries in the directories listed in the environment variable **OCPI_COMPONENT_LIBRARY_PATH**, which are colon separated and searched in order.
- The other component libraries in the same project.

- The component libraries exported by the projects listed in the environment variable `OCPI_PROJECT_PATH`, which are colon separated and searched in order.
- The component libraries of the projects listed in the `ProjectDependencies` variable in the project's `Project.mk` file.
- The component library in the *OpenCPI CDK*, identified in the environment variable `OCPI_CDK_DIR`.

In all cases, for a worker to be found, it must have been built.

5 HDL Primitives

HDL Primitives are HDL assets that are lower level than workers and may be used as building blocks for workers. HDL primitives can either be **libraries** or **cores**.

HDL primitives are useful for HDL workers when there is lower level code that is reused or shared in different workers. Using HDL primitives is also useful when there are non-OpenCPI code modules that are imported and should be left untouched in order to remain useful outside of OpenCPI. The use of HDL primitives is *not* required for HDL workers. HDL primitives cannot have the same name as a worker.

When lower level code modules and files are used in only a single worker, there is no need for a primitive library: such files can be simply put in the worker's directory and added to the **SourceFiles** variable in the worker's **Makefile**. Such files are built before the worker source files so that they can easily be referenced without forward declarations (e.g. **component** declarations in VHDL).

An **HDL Primitive Library** is a collection of low level modules compiled from source code that can be referenced in HDL worker code. An HDL worker declares which HDL primitive libraries it uses.

An **HDL Primitive Core** is a single low level module that may be:

- Built and/or synthesized from source code
- Imported as presynthesized and possibly encrypted, from a third party.
- Generated by tools such as Xilinx CoreGen or Altera MegaWizard.

An HDL worker declares which primitive cores it requires (and instantiates).

In both cases the exported/installed library or core that results from building a primitive is something that can be referenced by workers simply by including the following lines in the HDL worker **Makefile**:

```
Libraries=myutils
```

or

```
Cores=mycore
```

If the library or core name is a simple name (no slashes), then it is found by searching as described earlier; otherwise it is a path name to the primitive's directory.

When the worker source code instantiates a primitive core or a module from a primitive library, no further action needs to be taken other than including the line above in the HDL worker's **Makefile** (or once for all workers in a component library in the component library's **Makefile**). In particular, no other "black box" module or VHDL component declaration needs to be created by the worker.

The **HdLLibraries** variable can be set in the library's **Library.mk** file or the project's **Project.mk** file to make HDL primitive libraries available to all HDL workers in the library or project. The OpenCPI CDK itself includes several HDL primitive libraries, and some are always available for use by all workers, even when the **Libraries** variable is not set.

5.1 HDL Primitive Libraries

Primitive libraries are normally created using the `ocpidev` tool in a project, e.g.:

```
ocpidev create hdl primitive library myprims
```

This creates a directory for the primitive library in the project, in the directory `hdl/primitives`, with a `Makefile` in that directory containing:

```
include $(OCPI_CDK_DIR)/include/hdl/hdl-library.mk
```

This command can also be issued in the `hdl/primitives` directory itself.

Primitive libraries may be written in VHDL or Verilog, but there are specific rules to follow (mentioned above) in order for the library to be usable with all supported tools, and from VHDL or Verilog. Primitive libraries can depend on other primitive libraries, and this must be indicated by setting the `Libraries` variable in the `Makefile`, in the same way as it may be set in a worker's `Makefile`. Circular dependencies among primitive libraries are not supported. Some internal OpenCPI primitive libraries are always available to other primitive libraries libraries. This is suppressed if the `HdlNoLibraries` variable is set non-empty in the primitive library's `Makefile`.

5.1.1 Source Files in Primitive Libraries

If there are no ordering dependencies between source files, just creating or copying source files into the directory will cause them to be built there, together as the library. Thus without mentioning source file names, all source files in the top level directory of the primitive library will be built and included in the library.

The default build order for the source files in a primitive library directory is to first build any `*_pkg.vhd` and `*_body.vhd` files (see below) and then build all other source files (`*.vhd` and `*.v`). There are several conditions where all source files must be explicitly mentioned in the `SourceFiles` variable in the library's `Makefile`. These are:

- There are ordering dependencies between source files (other than dependencies on the `<pkg>.pkg.vhd` files which are always compiled first).
- Some source files are not in the top level directory of the library *and* are not in target-specific subdirectories for shadowing purposes (see [Target-Specific Modules](#) below).
- Some source source files that are in the library's directory should not be built into the library (i.e. extraneous unbuilt source files)

It is recommended that source files *not* be placed in subdirectories *unless they are there as target-specific modules*. The reasons this is not recommended are:

- there is then a potential name collision between the names of the subdirectories and the names of the HDL targets, tools, and vendor names used for target-specific modules
- all the files must be listed in the `SourceFiles` variable, which is otherwise usually unnecessary

All modules in a library intended to be used from outside the library must be in separate source files with the name of the file matching the name of the module, including case (before the language suffix). While not strictly required, this practice is also recommended for modules instantiated by other modules in the same library.

5.1.2 Package Declarations for Primitive Libraries

The library must include a VHDL file `<libname>.pkg.vhd`, containing component and data type declarations for all modules externally referenced (from outside the library). Even if the library has Verilog source code modules, the `<libname>.pkg.vhd` VHDL file with component declarations must be present for all of the Verilog modules in the library that are usable from outside the library.

The VHDL package name in the `<libname>.pkg.vhd` file should normally be the same as the library's name, but this is not required. There can be multiple `<pkg>.pkg.vhd` files in a library if multiple packages are required. Finally, if the packages have package bodies in separate files, those files should be named `<pkg>-body.vhd`.

When modules in the library are instantiated by other modules in the library, but *not* intended for external usage by code *outside* the library, their still-required component declarations can be placed in a different package to keep them separate from those in the `<libname>.pkg.vhd` file that *are* intended for external use. Perhaps such a package would be called `<libname>.internal_pkg.vhd`.

For example, consider the source file implementing module `outer`, in the file `outer.vhd`, in the `mylib` primitive library.

```
entity outer is
    port (clk : in std_logic; ...);
end entity outer;
architecture rtl of outer is begin
    fifo: component myownfifo port map(...);
end rtl;
```

For this module to be usable from outside the library, the component declaration must be in the library's package file, `mylib_pkg.vhd`:

```
package mylib is
    component outer is
        port (clk : in std_logic; ...);
    end component outer;
end package mylib;
```

In the example above, the `outer` module uses another module, `myownfifo`, also from this library. Any module in the library referenced between files, must also have a component declaration in a package file, so the package file might in fact be:

```

package mylib is
    component outer is
        port (clk : in std_logic; ...);
    end component outer;
    component myownfifo is
        port (...);
    end component myownfifo;
end package mylib;

```

Alternatively an internal package file could be defined, in `mylib_internal_pkg.vhd`, containing just the internally-used module:

```

package mylib_internal is
    component myownfifo is
        port (...);
    end component myownfifo;
end package mylib;

```

5.1.3 Instantiating Modules in Primitive Libraries

Modules referenced from VHDL must use the component instantiation syntax, so the instantiation does that in the `outer` entity example above.

When a VHDL worker or code in another primitive library is written to *use* a module in a primitive library, it must include a line to access the library. For a primitive library `mylib`, the calling module file might contain the line.

```

library mylib; use mylib.mylib.all;
...
inst1 : component outer ...

```

If there were multiple packages in the primitive library, the package names might be different than the library name, e.g.:

```
library mylib; use mylib.mypkg.all;
```

Alternatively, to avoid any name collisions but be more verbose, the code could be:

```

library mylib;
...
inst1 : component mylib.mylib.outer ...

```

Finally, when instantiating a module from within the same library, the work library can be used and no library declaration is required, e.g.:

```
inst1 : component work.mylib.outer ...
```

5.1.4 Providing Target-specific or Vendor-specific Versions of Primitive Modules

The modules in a primitive library are normally each in their own files in the top level directory with the file name being the same as the module name. Sometimes it is useful to have special versions of source code for a module that is specific to a particular part family, vendor or tool. This allows alternative source code files for a module that uses vendor-specific primitives or are written in a way to infer vendor-specific hardware features (e.g. BRAMs, DSPs, IO features).

A primitive library is normally built using each source file indicated in the `SourceFiles` variable, or if that variable is not specified, each source file found in the primitive library's directory. However, when building for a specific target, it first looks for a file of the same name in a subdirectory with the name of the target being built (e.g. `zynq` or `isim`). If it finds that file, it uses it *instead* of the file in the top level directory. If there is no such file in a target-specific directory, it next looks in a vendor-specific directory (e.g. `xilinx` or `altera`). If the file does not exist in either target-specific or vendor-specific subdirectories, the (default, generic) file in the top level directory is used.

Note that simulator targets also have vendors. Thus if the target is `isim` or `xsim` (both `xilinx` simulators), and there is a module file in the `xilinx` subdirectory, that file will be used for those simulators in preference to the default file at the top level.

The files in target-specific or vendor-specific directories should never be mentioned in the `SourceFiles` variable. The file in the top level directory serves as the default implementation of the module, which will be ignored (shadowed) in preference to target-specific or vendor-specific versions when they exist.

Two examples of this feature are below.

5.1.4.1 Shadowing Example for Correct Inference of Resources

A synchronous ROM module is defined to take CLK and ADDR as input, and provides output DATA synchronously after two clock edges, based on the address valid at that time, with no overlap of access cycles. A simply default (Verilog) implementation, in `ROM.v`, assuming single-cycle access times, would simply ignore the clock and drive data continuously:

```
assign DATA = ROM[ADDR] ;
```

This is correct functionality, perhaps suitable for simulation, but neither Xilinx nor Altera tools will synthesize this into their respective block ram resources. For Xilinx, the output data must be registered for the synthesis tool to infer/use a block ram, thus the Xilinx code should be:

```
always @ (posedge CLK) begin
    DO_R <= ROM[ADDR] ;
end
assign DO = DO_R;
```

For Altera, the input address must also be registered, thus the code should be:

```
always @ (posedge CLK) begin
    ADDR_R <= ADDR;
    DO_R <= ROM[ADDR_R] ;
end
assign DO = DO_R;
```

The contract of the module allows for all three implementations:

- `ROM.v` for a simple default implementation suitable for simulation.
- `xilinx/ROM.v` to correctly infer BRAMs using ISE tools.
- `altera/ROM.v` to correctly infer BRAMs using Quartus tools.

5.1.4.2 Example of Shadowing using Vendor-specific Primitives

A clock buffer is an important resource for clock distribution, and it is desirable to write higher level code that uses portable primitives to instance one. This example takes CLK as input and produces CLK_BUFFERED as output. The easiest way to use clock buffer resources is to directly instance the vendor-specific primitives. A generic clock buffer default implementation, in `clkbuffer.vhd`, would be:

```
clk_buffered <= clk;
```

This is functionally correct, but does not take advantage of clock buffering or routing features unless recognized automatically due to fanout etc. For Xilinx, the explicit implementation would use a BUFG primitive, e.g.:

```
buf : BUFG port map(I => clk, O => clk_buffered);
```

For Altera, an attribute declaration might be sufficient for this purpose:

```
attribute altera_attribute of clk_buffered :  
    signal is "-name GLOBAL_SIGNAL REGIONAL_CLOCK";
```

The module would have three implementations:

- `CLK_BUFFERED.vhd` for a default/simple/portable implementation for simulation
- `xilinx/CLK_BUFFERED.vhd` to instance BUFG primitive for Xilinx
- `altera/CLK_BUFFERED.vhd` to use an explicit Altera attributes

5.1.5 Exporting and Using the Results of Building HDL Primitive Libraries

When HDL primitive libraries are built, their immediate per-target results are in target-specific subdirectories (`target-<hdl-target>`) whose format varies depending on the tools used for that target. The log output of the tools is usually collected in a `<libname>-<tool>.out` file in the target directory, which can be examined when errors occur or to examine warnings, etc.

HDL Primitives are always built as a group under the `hdl/primitives` directory in a project, with its own `Makefile`, which is automatically created whenever primitives are created using `ocpidev` in a project.

Building a primitive library in its own directory is useful for rapidly getting to a clean build across all relevant targets. To make these results available to workers and other projects, the primitives in an `hdl/primitives` directory (in a project), must be built from a higher level, even if done individually. When that is done, the exportable results are placed in `primitives/lib`, much like the `lib` subdirectory of component libraries. The files in that directory are automatically used as the project's exported primitive libraries and cores when a project is built. Building a project from the top level builds and exports all the primitives automatically.

When developing a primitive library it is highly recommended to at least occasionally build for all available tools, both for synthesis to hardware and for simulation. This ensures that the code is nominally portable.

5.2 HDL Primitive Cores

Making a prebuilt/presynthesized core available for use by workers is similar to creating a primitive library from source files. The `ocpidev` command, for creating `mycore`, is:

```
ocpidev create hdl primitive core mycore
```

This creates a directory for the primitive core in the project, in the directory `hdl/primitives`, with a `Makefile` in that directory containing:

```
include $(OCPI_CDK_DIR)/include/hdl/hdl-core.mk
```

This command can also be issued in the `hdl/primitives` directory itself.

Whereas an HDL primitive library is built as a collection of source modules that are not fully elaborated or synthesized, HDL primitive cores are built into a *single* module that may have no source files other than those that define the interface.

The files used to build the core can be a mix of source and prebuilt files. There may be presynthesized core files (e.g. Xilinx `.ngc` or Altera `.qxp`), or source files. There may be presynthesized files for some targets and source files for other targets. Any targets that do not have prebuilt cores will use the source files.

When the core supports instantiation from Verilog, there must be a “black box” empty module definition file `<corename>_bb.v`. When the core supports instantiation from VHDL, it must have a package file `<corename>_pkg.vhd` containing a component definition in a package named the same as the core name.

There are two special additional optional `Makefile` variables that apply to HDL primitive cores: `Top` and `PrebuiltCore`.

`Top` is the variable that specifies the top module name of a primitive core when it is different from the core name used when it was created with `ocpidev`. Normally the name of the primitive core is the same as the top level module name and this variable is unnecessary. In some cases the core name is more descriptive and useful, while the top module name is predetermined for some other reason. An example is a core name of `ddc_4ch_v5`, which might be a core for a 4 channel DDC generated specifically for virtex5. The actual generated core from Xilinx CoreGen has the file and module name `duc_ddc_compiler_v1_0`. Thus the `Makefile`, in the `ddc_4ch_v5` directory would contain:

```
Top=duc_ddc_compiler_v1_0
```

The `PrebuiltCore` variable is used to specify a file that is a core that is not in source code, but is generated by some other tool and copied into the directory for the HDL primitive core. If this variable is not set, source files are expected. Here is an example:

```
PreBuiltCore=mycore.ngc # suppress building from source files
OnlyTargets=xcv6lx240t # this core is only good for this part
include $(OCPI_CDK_DIR)/include/hdl/hdl-core.mk
```

In cases where there are different prebuilt core files for different targets, the variable name `PrebuiltCore_<target>` is used. Thus if a core named `fft4k` supported

both VHDL and Verilog, and had prebuilt cores for zynq and stratix4, as well as source code for other (typically simulation) targets, the **Makefile** could contain:

```
PrebuiltCore_zynq=myzynqfft4k.ngc
PrebuiltCore_stratix4=mys4fft4k.qxp
SourceFiles=fft4kforsim.v
```

The directory for the primitive core would contain these files:

```
Makefile
fft4k_bb.v
fft4k_pkg.vhd
myzynqfft4k.ngc
mys4fft4k.qxp
fft4kforsim.v
```

If the source files were specific to the Xilinx Isim simulator, then the core should be restricted to building only for the zynq, stratix4, and Isim targets using, in the **Makefile**:

```
OnlyTargets=isim stratix4 zynq
```

Primitive cores can depend on other primitive cores or libraries, and this must be indicated by setting the **Libraries** or **Cores** variable in the **Makefile**, in the same way as it may be set in a worker's **Makefile**. Circular dependencies are not supported.

6 HDL Assemblies for Creating Bitstreams/Executables

An HDL assembly is a fixed composition of HDL application workers that can act as a whole or part of a heterogeneous OpenCPI application. It is built into multiple FPGA bitstreams for different FPGA platforms. It will execute as part of some OpenCPI application with its workers being a subset of the workers in the application.

This section describes how to define and build (synthesize) these assemblies and to ultimately turn them into bitstreams. When the target HDL platform is a simulator, we use the term **executable** while when the target is an actual physical FPGA, we use the term **bitstream**. In both cases there is a single resulting standalone artifact file that is ready for loading and execution.

This process *implements the assembly onto platforms* as part of an application. More specifically, it is combining the defined assembly of HDL workers with a **platform configuration** to create an **HDL container** that is then transformed into a bitstream/executable. This was shown in the build hierarchy in the section [HDL Build Hierarchy](#).

The container is the outer module that contains the HDL assembly as well as the platform support modules in the platform configuration. A platform configuration is for either a simulation platform or a real FPGA platform.

Given that the platform configurations already exist, the steps taken to go from the assembly (described in XML) to a bitstream are:

1. Describe the assembly in XML, specifying application workers and connections.
2. Select the platform configuration that the assembly will be implemented on.
3. Specify how the assembly's external ports connect to the platform (i.e. to an interconnect like PCIe for off-platform connections, or to local devices). This is "defining the container".
4. Run **make** to generate the bitstream.

In most cases, steps #2 and #3 above are automatic and use defaults. A quick example, typically used for unit testing would be an assembly file containing a single worker:

```
<HdLAssembly>
  <Instance Worker="bias_vhdl" externals='true' />
</HdLAssembly>
```

HDL assemblies are created using the **ocpidev** tool, e.g. for creating the **myassy** assembly:

```
ocpidev create hdl assembly myassy
```

This creates a directory in the project's **hdl/assemblies** directory, with the name of the assembly (**myassy**), containing an initial HDL assembly XML file and an initial **Makefile**. The name of the XML file is simply the name of the assembly, with the ".xml" file extension. This command can also be issued in the project directory or the **hdl/assemblies** directory of a project. It can also be run standalone, using the **ocpidev**

-s option, in any directory. When the command is issued for the first time in a project, the **hdl/assemblies** directory itself will be created, with its own **Makefile**. Thus after creating the first HDL assembly (called **first**) in a project, and building it for the **zed** platform, the directory structure would be:

```
hdl/assemblies/          # Directory for all assemblies
    Makefile            # Makefile for all assemblies
    first/              # Directory for first assembly
        Makefile          # Makefile for the first assembly
        first.xml          # XML file for the first assembly
    -- from here down is results from building for zed --
    gen/xyz-assy.v       # generated assy structural hdl
        -- other generated files --
    target-zynq/         # synthesized assembly build dir
    container-first_zed_base/ # container dir for first on zed
        gen/first_zed_base-assy.vhd # generated top level vhdl
        -- other generated files --
    target-zynq/         # final bitstream build directory
        first_zed_base.bitz      # final bitstream/executable file
```

The actual steps taken by the OpenCPI scripts and tools, to create a bitstream or executable from an assembly, are:

1. Generate the Verilog/VHDL code that structurally implements the assembly.
2. Build/synthesize the assembly module, that has some “external ports”.
3. Generate the Verilog/VHDL container code that structurally combines the assembly and the platform configuration, as well as any necessary device workers.
4. Build/synthesize the container code, incorporating the assembly and platform configuration. This is the top level module.
5. Run the final tool steps to build the bitstream (map, place, route, etc.).

HDL adapters will automatically inserted as needed in both steps 1 and 3.

This process is run for all platforms specified in the **Hd1Platforms** variable, as specified on the command line or in the project's **Project.mk** file.

This results in an artifact file, with the suffix **.bitz**, which can be used at runtime when executing OpenCPI applications. This file is based on the vendor-tool-specific output files like **.bit** for Xilinx and **.sof** for Altera, which are also in the container/bitstream build directory. Those files are not used by OpenCPI after the **.bitz** file is created.

6.1 The HDL Assembly XML file

The assembly is described in an XML file containing an `HdlAssembly` top-level XML element, which contains **worker instances**, **property/parameter settings**, **connections** and **external ports**. It is similar to the Application XML file that describes the whole OpenCPI heterogeneous application (as documented in the [OpenCPI Application Guide](#)).

The XML file is generated initially when the `ocpidev create hdl assembly` command is issued.

The worker instances (`instance` subelements of the `HdlAssembly`) reference HDL workers in some component library, and optionally assign names to each instance. The `worker` attribute is the worker's OWD name (without directory or model suffix), and the optional `name` attribute is the instance name. When not specified, instance names are either the same as the worker name (when there is only one instance of that worker in the assembly), or the worker name followed directly by a zero-based decimal ordinal (when there is more than one instance of the same worker).

Connections among workers in the assembly can use `connection` XML elements or a more compact shorthand described next. The `connection` elements define connections among worker data interfaces. A trivial example would be:

```
<HdlAssembly>
  <Instance Worker="generate"/>
  <Instance Worker="capture"/>
  <Connection>
    <port name='out' instance='generate' />
    <port name='in' instance='capture' />
  </Connection>
</HdlAssembly>
```

For convenience, internal connections between the output of one instance to the input of another can simply be expressed using the `connect` attribute of the instance, indicating that the instance's only output should be connected to the only input of other instance whose name is the value of the `connect` attribute. The example above can be written as:

```
<HdlAssembly>
  <Instance Worker="generate" connect='capture' />
  <Instance Worker="capture"/>
</HdlAssembly>
```

Furthermore, when this shortcut is used, you can specify the “from” port using the `from` attribute, and the “to” port using the `to` attribute. If these instances had multiple other input and output ports, you can also specify it this way:

```
<HdlAssembly>
  <Instance Worker="generate" connect='capture' from='out' to='in' />
  <Instance Worker="capture"/>
</HdlAssembly>
```

To specify external ports, where the data is flowing into or out of the assembly itself, the **external** element is used, which allows the name of the external port to be different from the worker port it is connected to:

```
<Hd1Assembly>
  <Instance Worker="generate" connect='process' from='out' to='in' />
  <Instance Worker="process"/>
  <external name='procout' instance='process' port='out' />
</Hd1Assembly>
```

But there is also a shortcut when the external port name is the same as the worker's port name, by simply using the **external** attribute of the instance:

```
<Hd1Assembly>
  <Instance Worker="generate" connect='process' from='out' to='in' />
  <Instance Worker="process" external='out' />
</Hd1Assembly>
```

To specify that all unconnected ports of a worker should be made external ports of the assembly, you can use the **externals** boolean attribute. E.g., if the assembly is in fact a single worker where both **in** and **out** ports should be external, you would only need:

```
<Hd1Assembly>
  <Instance Worker="process" externals='true' />
</Hd1Assembly>
```

This is a common use case for unit test assemblies used to test single HDL workers.

Below is a diagram of a simple assembly, and the corresponding Hd1Assembly XML file. The application has a **switch** worker that accepts data either from its **in0** or **in1** interface, and sends the data to its **out** interface. The **delay** worker sends data from **in** to **out** implementing a delay-line function that requires memory. The **split** worker takes data from its **in** interface and replicates it to both its **out0** interface as well as its **out1** interface. The HDL assembly has 4 external ports (ADC, SWIN, SWOUT, DAC).




Figure 4: Example HDL Assembly

Given that these three workers are already in a component library, the XML description of the example is below. It uses the **external** elements rather than simply adding the

`externals='true'` to the switch and split workers because the external ports of the assembly have different names than the corresponding worker ports.

```
<HdlAssembly>
  <Instance Worker="switch" connect="delay"/>
  <Instance Worker="delay" connect="split"/>
  <Instance Worker="split"/>
  <External name='adc' instance="switch" port="in0"/>
  <External name='swin' instance="switch" port="in1"/>
  <External name='dac' instance="split" port="out0"/>
  <External name='swout' instance="split" port="out1"/>
</HdlAssembly>
```

Figure 5: Example HDL Assembly XML

6.2 The **Makefile** for Building an Assembly

The HDL assembly **Makefile** is automatically created when the `ocpidev create hdl assembly` command is issued. The **Makefile** indicates which component libraries should be used to find the workers mentioned in the **Hd1Assembly**. A **Makefile** for the above application might be:

```
ComponentLibraries=/home/fred/project/components
include $(OCPI_CDK_DIR)/include/hdl/hdl-assembly.mk
```

In this case the assembly's workers will be found in the indicated component library, or other libraries in the search path, described earlier in the [HDL Search Paths when Building](#) section. This variable can also be set in the **hdl/assemblies** directory **Makefile** to apply to all the assemblies in that directory (see just below), or in the project's **Project.mk** file.

While HDL primitives and workers are built for FPGA part families, indicated using **Hd1Targets**, HDL assemblies are built for specific HDL *platforms*, using the **Hd1Platforms** variable. The platform includes specific devices attached to the FPGA as well as other specific FPGA attributes. Anywhere that **Hd1Targets** is required, **Hd1Platforms** can be used, since a platform implies a target. The variables suitable for HDL assembly MakeFiles are:

*Table 20: HDL Assembly **Makefile** Variables*

Variable Name in HDL Worker Makefile	Usable as default in hdl/assemblies Makefile ?	Description
ComponentLibraries	Y	A list of component libraries to search for the workers in the HDL assembly (in order)
OnlyPlatforms	Y	An exclusive list of platforms for which this assembly (and default containers) should be built.
ExcludePlatforms	Y	A list of platforms for this the HDL assembly should <i>not</i> be built.

When you have a directory full of HDL assemblies like this, they are usually in the **hdl/assemblies** directory of a project, acting as a sort of library of assemblies. This **Makefile** in the **hdl/assemblies** directory is automatically created when the first HDL assembly is created. It can simply be:

```
include $(OCPI_CDK_DIR)/include/hdl/hdl-assemblies.mk
```

If the assemblies in this directory should not all be built, or should be built in a particular order, the **Assemblies** variable can be set to a list of assemblies to be built in the specified order.

6.3 Specifying the Containers that Implement the Assembly on Platforms

The container is the top-level module and implements the assembly on the platform. Separating the assembly from the container in this way keeps the assembly portable and hardware-independent (assuming the workers are). The assembly's external input and output connections are unspecified until it is implemented in a container on the platform. By itself, the HDL assembly is usable in a simulation testbench, when it is built for a simulation platform.

In an assembly's **Makefile**, containers are specified in two ways. The first is **default containers**. Default containers are generated by looking at the assembly, and connecting all the external ports to the platform's interconnect (e.g. PCI Express or SoC buses). Thus the implicit default container specification is to connect every external port of the assembly such that it connects external to the platform, to connect to workers running on containers on other platforms. Using the previous assembly example, the default container would implement the assembly with all 4 ports (ADC, SWIN, DAC, SWOUT) connected to the platform's interconnect.

The optional **DefaultContainers** variable is used to list platform configurations for which default containers (and thus bitstreams) should be automatically generated for this assembly. The format of the items in this list is **<platform>** or **<platform>/<configuration>**. If this variable is defined as empty:

```
DefaultContainers=
```

no default containers are built for this assembly. If the **DefaultContainers** variable is not set at all (not mentioned in the assembly's **Makefile**, the default situation), then default containers (and bitstreams) will be generated for whatever platform the assembly is built for. In this case the platform configuration is assumed to be the **base** platform configuration for the platform (the one with no device workers at all).

Based on these defaults, if nothing is said at all about containers in the **Makefile**, default container bitstreams will be built for whatever platforms are mentioned in **Hd1Platforms**. In many cases this default case is all that is required (no container variables at all).

Default containers are used to connect the external ports of the assembly to the single system interconnect of the platform. I.e. if there are multiple interconnects (say PCIe and Ethernet), or if connections to local devices are required, then a container must be specified in its own XML file. The **Containers** variable specifies a list of containers (container XML files) that should be built in addition to those indicated by the **DefaultContainers** variable (or absence thereof). The **Containers** variable does not suppress the building of the default containers.

This example **Makefile** relies only on the component libraries specified in search paths, or in the **ComponentLibraries** variable setting in the **hd1/assemblies/Makefile**, and builds default containers for whatever platform is specified. It is essentially what is generated automatically by the **ocpidev** command.

```
include $(OCPI_CDK_DIR)/include/hd1/hd1-assembly.mk
```

A **Makefile** that builds a default container on the **m1605** base platform configuration and the **lime_adc** configuration of the **alst4** platform, and further generates a specific container called **in_2_adc** for connecting some external port to the ADC device on that latter platform configuration might look like:

```
DefaultContainers=m1605 alst4/lime_adc
Containers=in_2_adc
include $(OCPI_CDK_DIR)/include/hdl/hdl-assembly.mk
```

6.4 HDL Container XML files

A container XML file is required to connect to multiple interconnects (or not the first one) or to make connections to local devices.

The top-level element of the container XML file is the `Hd1Container` element. Its primary attribute is the `platform` attribute to specify the platform configuration that should be targeted by the container. It is of the form `<platform>` (which implies the base configuration) or `<platform>/<configuration>`. Beyond specifying the platform configuration in the `platform` attribute, the subelements of the `Hd1Container` top-level element are `connection` elements for containers, with these attributes:

- `external`: specify an external port of the assembly
- `device`: specify a device on the platform or on a card
- `interconnect`: specify an interconnect on the platform
- `port`: specify the device's port (required when device has more than one data port)
- `otherdevice`: specify a second device for device-to-device connections
- `otherport`: specify the otherdevice's port for device-to-device connections
- `card`: specify a type of card that a device is on (required if device is not part of the platform, i.e. not defined in the platform's XML)
- `slot`: specify the slot that a card is plugged into for a device (required when the card is supported in more than one of the platform's slots)

Using these attributes you can specify connections between:

- external and interconnect
- external and device
- interconnect and device
- device and otherdevice

Device-to-device connections are currently only supported when neither is on a card, or both are on the same card.

In a case where the `in` of the assembly connects to a locally attached `adc` device, but the `out` of the assembly is attached to the interconnect for communicating to other FPGAs or software containers, you would have:

```
<Hd1Container platform='alst4/alst4_conf1'>
  <connection external='in' device='adc' />
  <connection external='out' interconnect='pcie' />
</Hd1Container>
```

The connection to the `adc` device will be resolved in one of two ways:

- if the `adc` device was already included in the `a1st4_conf1` configuration of the `a1st4` platform, then that `adc` device instance will be used
- If not, then the `adc` device logic would be instantiated in the container itself, and used there

6.4.1 Service Connections in a Container

OpenCPI HDL workers have three types of ports: control, data, and service. Service ports are connected locally in the container to provide the required services to workers. Service ports are implementation-specific for a given worker and thus not found in OCS files. I.e., the worker declares what services it needs for its particular implementation. The services currently defined are memory and time (time of day).

If workers in the assembly have service ports, they automatically become external ports of the assembly, but not for data. When generating a container, all services required by the assembly, as well as any services required by device workers instanced in the container, must be satisfied by instancing the appropriate service modules in the generated container code.

Time service requirements (based on `timeinterface` elements in the worker's OWD) are satisfied by:

- instancing a ***time client*** module for each worker that needs “time of day”, and
- connecting that ***time client*** to the timekeeping infrastructure on the platform

Each time client is instanced based on requirements of the associated worker’s time port, as specified by the `timeinterface` element in the worker’s OWD.

Memory service requirements [which are currently not supported as of this writing] are satisfied by instancing either private BRAM modules (private to the worker) or instancing external memory access interfaces connected to device workers for external memory. Memory access may also be multiplexed to support multiple workers sharing the same memory. Note that such memory services are unrelated to data message buffering used to connect workers together or connect them to devices and interconnects.




Figure 6: OpenCPI HDL Assembly on Container and Platform

6.4.2 Preparing the Bitstream/Executable Artifact File

This section describes how the container build process injects metadata into the resulting artifact file. There is no user control over this process, but it is useful to understand for troubleshooting purposes. Containers are built into bitstream/executable files by the FPGA back-end place-and-route or simulation tools. These tool-specific files are then post-processed into a generic `.bitz` file that acts as an OpenCPI artifact for application execution. This post-processing always compresses the tool-specific output files using `gzip`.

This post processing also uses an ***artifact XML*** file that is generated and describes what is in the bitstream/executable. This information includes the contents of the platform configuration, the assembly, and the container. The post-processing attaches ***artifact XML*** to the file in two ways.

The first is that the XML is compressed and embedded into the logic of the bitstream in a block memory. This allows OpenCPI software to extract it when the bitstream is loaded in an FPGA. It enables software to know what is in the bitstream without knowing or having access to the file it was originally loaded from. E.g., if the FPGA was booted from a flash memory attached directly to the FPGA and not accessible to software, software can still retrieve this information and know how to use the bitstream. The `ocpihdl` tool, described in the [***ocpihdl command-line utility***](#) section, is used to manually query this embedded information for troubleshooting purposes.

The second way the XML is attached to the file is outside the FPGA configuration logic, but attached to the bitstream (`.bitz`) file so that it can be retrieved from the file

generically, regardless of the FPGA or simulation tools used to create the file. This allows software to know what is in the file, without it being loaded into a device, and without knowing any other files or locations that the artifact file came from. It makes the file self-describing. This attached XML is the same for all artifact files for all authoring models. The `ocpixml` command-line tool is used to extract this XML data into a separate file for examination.

In summary:

- target-specific FPGA or simulation tools create the raw container output file
- this file has embedded artifact XML in an initialized block memory
- this file is compressed
- the artifact XML is then attached to the compressed file like all other artifact files

Artifact files are installed in an OpenCPI runtime component library as referenced by the OCPI_LIBRARY_PATH environment variable. It is the format expected by the internal OpenCPI mechanisms to load bitstreams onto platforms at runtime.

To conveniently collect all the bitstream/executable artifact files in a project in one place, they can be exported putting the following line in the `Project.exports` file:

```
+hdl/assemblies/*/container-*/*target-*/*.bitz lib/hdl/assemblies/
```

Thus all the bitstream/executable artifact files in the project will be accessible in the `exports/lib/hdl/assemblies` directory of the project, and that directory can be placed in the OCPI_LIBRARY_PATH environment variable that is used to find artifacts when executing OpenCPI applications.

7 HDL Simulation Platforms

The OpenCPI concept of an HDL platform encompasses both physical FPGA-based platforms as well as HDL simulators such as Mentor's modelsim, and Xilinx isim.

At build time, simulators are “just another target” when building primitives and workers, and “just another platform” when building assemblies and containers. The names for simulators in both target-related and platform-related `make` variables are `modelsim` and `isim`. E.g.:

```
make Hd1Platforms="zed alst4 modelsim"
```

would build for the ZedBoard Zynq-based platform, the Altera Stratix4-based platform, and the modelsim platform. If a worker was intended only for simulators, its `Makefile` would typically contain:

```
OnlyTargets=isim modelsim
```

Similarly, if an assembly was intended only for simulators, its `Makefile` would typically contain:

```
OnlyPlatforms=isim modelsim
```

Even though simulators do not perform synthesis, building for simulators does what is possible, and tries to elaborate the design at each build level to catch errors as early as possible. To suppress this incremental elaboration for faster build times, the `Makefile` variable `Hd1NoSimElaboration` can be set to 1. Since simulation builds are normally fast anyway, this is rarely worth it, especially for modelsim.

At runtime, an installed simulator is an available HDL platform just like any installed hardware HDL platform. Simulators act as much like a hardware FPGA platform as possible:

- It is discoverable (using `ocpihdl search`, or `ocpirun -C`)
- It appears as available without a bitstream being loaded

Previous versions of OpenCPI used the `ocpihdl simulate` command to start a simulator container as a server process. This is no longer necessary or supported. For backward compatibility for some test benches, this command does nothing but sleep indefinitely.

The following simulation server behavior is currently *disabled*, but will be enabled in a future release.

- A bitstream can be “loaded”, which in fact starts simulation (usually automatically as needed by `ocpirun`, or explicitly using `ocpihdl load`)
- It is persistent, so an application can be executed multiple times, which in fact will occur in the same simulation run
- It can be queried (e.g. to find the current value of a worker’s properties) using various `ocpihdl` commands

7.1 Execution of Simulation Bitstreams and Containers.

Simulators that are installed for use by OpenCPI are automatically available as containers, similar to HDL hardware platforms or even software containers. When an application is run, these containers are available to be used if artifacts are built and accessible via the `OCPI_LIBRARY_PATH` environment variable. For example, assuming that modelsim and isim simulators are installed for OpenCPI, on a CentOS7 system, the command `ocpirun -c` would output:

```
Available containers:
# Model Platform      OS     OS-Version Arch      Name
0 hdl    isim          -       -           -        lsim:isim
1 hdl    modelsim      -       -           -        lsim:modelsim
2 rcc    centos7       linux  c7         x86_64   rcc0
```

If an application is run, and bitstream files are available to use these simulators, they will be used automatically. To force components in the application to use a particular simulator, the `-P` option to `ocpirun` can be used, e.g.:

```
% ocpirun -P=modelsim myapp
```

would run `myapp`, forcing all components to be executed with modelsim. To force one component in the application to use modelsim, you could say:

```
% ocpirun -Pmycomp=modelsim myapp
```

The `ocpirun` command normally used to run OpenCPI applications has several options that apply only to simulators:

Table 21: Simulations Options to `ocpirun`

Name	Letter	Description
<code>sim_dir</code>	<code>none</code>	The name of a directory where simulation outputs will be placed. The default is <code>simulations</code> , relative to where <code>ocpirun</code> itself is running.
<code>sim-ticks</code>	<code>none</code>	The number of simulation clock cycles to execute or until the application is done.

A complete description of the `ocpirun` command is in the [*OpenCPI Application Development Guide*](#).

As with execution on any hardware HDL device (FPGA), some of the components in the application may be in software containers running on the host processor. The data connections between workers inside the HDL device (or simulator) and outside the HDL device in software containers work normally. Since HDL execution is much slower in simulators than in real FPGA hardware, the software workers will see data consumed or produced by the HDL simulator device much slower.

When some OpenCPI application (e.g. executed using the `ocpirun` utility command) decides to run an assembly of workers on this simulator-based device (as it would with any other discovered and available HDL device/FPGA), it would request that the

bitstream (executable) be “loaded” and “started” on this device. This would cause this simulated HDL device to run the actual simulator (e.g. modelsim) with that executable.

Each time a simulator is actually run under `ocpirun`, it will execute in a new subdirectory created for that simulation run, with the name:

`<assembly-name>.<sim-platform>.<date-time>`

Thus running applications that use simulators will result in one or more subdirectories holding simulation results for each simulation run. No subdirectories are created when the simulator is simply discovered using `ocpirun -C` or `ocpihdl search`.

Each simulation execution that is launched by `ocpirun` when it runs an application will continue until one of the following occurs:

- The code being simulated explicitly asks for the simulation to terminate via the `$finish` system task in Verilog or an `assert` in VHDL.
- The simulation run exceeds the control-plane clock cycle count provided by the `--sim-ticks` option to `ocpirun`.
- This `ocpirun` command receives a control-C.

The results of any simulation run can be viewed using the `ocpiview` command. This command with no arguments opens the most recent simulation run found in the `simulations` directory, using the simulation viewer associated with the simulator used in that run. If given an argument, it is the directory containing a particular simulation run. The normal pattern of development is to run `ocpiview` after execution if examining the simulation run in detail is needed.

For Xilinx “isim”, the actual underlying viewing command (in the subdirectory for the simulation run) would be:

```
isimgui -view sim.wdb
```

For modelsim, it would be:

```
vsim -view vsim.wlf
```

In all cases the log of the simulator's output for the run is in the `sim.out` file.

8 HDL Device Naming

The term **HDL Device** is used here to refer to an instance of an HDL platform in an OpenCPI system (and not a device attached to an FPGA *inside* the platform). HDL devices have unique names within the *system*. HDL devices host HDL containers for execution. Each name starts with a prefix indicating how the device is discovered and controlled by OpenCPI software.

The control schemes currently supported are:

PCI

FPGA devices/boards accessible by PCI Express

Ether

FPGA devices accessible via link-layer Ethernet

LSim

FPGA devices that are in fact simulators (see the [**HDL simulation platforms**](#) section)

UDP

FPGA devices that are accessible via IP/UDP

PL

FPGA device in a Zynq SoC accessible via the on-chip AXI interconnect

The full device name is of the form:

`<control-scheme>:<address-for-control-scheme>`

As a convenience, if there is no known prefix in the name, then if there are 5 colons in the device name, it is assumed to be an **Ether** device name. Otherwise it is assumed to be a **PCI** device name.

8.1 PCI-based HDL devices

HDL platform devices on the PCI express bus/fabric are identified by the syntax common to many PCI utilities such as **lspci** on Linux, namely:

```
<domain>:<bus>:<slot>.<function>
```

An example is:

```
0000:05:00.0
```

Since it is common to have “domain”, “slot”, and “function” all being zero, if the address field in the device is simply a number with no colons, it is assumed to be the bus number with the other fields being zero. Thus “pci:5” implies “pci:0000:05:00.0”. Since an identifier with no prefix and not having 5 colons is assumed to be a PCI device, the identifier “4”, is assumed to be “pci:0000:04:00.0”.

A common example of a PCIe device is the Xilinx ML605 development board. Another is the Altera Stratix4 development board (called **a1st4** in OpenCPI).

8.2 Ethernet-based HDL Devices

HDL devices that are attached to Ethernet and operate at the link (or MAC) layer (OSI layer 2) use the **ether** prefix. This prefix implies access without using any routing or transport protocols such as IP/UDP or IP/TCP. It is the fastest and lowest latency way to use a network, with the drawback that it cannot be “routed” through IP routers, but can only be “switched” by L2 Ethernet bridges and switches. The syntax for naming such devices is:

```
Ether: [<interface>/]<mac-address>
```

The MAC-address is the typical 6 hexadecimal bytes separated by colons, such as:

```
c8:2a:14:28:61:86
```

The optional <interface> value is the name of a network interface on the computer accessing the device. Typical examples are **en0** or **eth0**. Newer Linux systems use more complex (but predictable) names that relate to busses and slots, e.g. **enp14s0** for PCI-based network interfaces. It is optional when there is only one such device. On systems with multiple interfaces it indicates which one should be used to reach the device. This is needed since there is no routing at this level of the network stack: you must use the right network interface to reach the addressed device.

The available network interfaces can usually be identified by the **ifconfig** Linux command. There is a more special purpose sub-command of the **ocpihdl** utility that lists only the network interfaces available and usable for OpenCPI (the **ethers** command to **ocpihdl**).

8.3 Simulator Device Naming

OpenCPI runs HDL simulators in a way that makes them look like any other device to software. When they are available (installed for OpenCPI), they are discoverable like any other device. They are accessed by name according to this syntax:

```
LSim:<simulator>
```

Current simulators are `modelsim` and `isim`. Such simulators are discovered automatically, and can be listed with the `ocpihdl search` command, along with hardware FPGA platform devices.

9 The **ocpihdl** Command Line Utility for HDL Development

The **ocpihdl** utility program performs a variety of useful functions for OpenCPI HDL development. These include:

- Searching for available FPGA devices (via PCI, Ethernet, UDP, simulators, etc.)
- Testing the existence of a specific FPGA device
- Reading and writing specific registers in an FPGA device
- Loading bitstreams on a device
- Extracting the XML metadata from a running device

The general syntax of ocpihdl is:

```
ocpihdl [<options>] <command> [<options>] [<command arguments>]
```

Options are the typical hyphen-letter options, some with arguments after the hyphen-letter argument. The following sections describe each command and its associated options and arguments. Here are the options that apply to many commands:

- v be verbose – put progress messages on standard output
- d **<hdl-device>**
specify the HDL platform device; see [HDL Device Naming](#) above
- p **<hdl-platform>**
specify the HDL platform type (e.g. ml605) for the command
- l **<log-level>**
specify the OpenCPI logging level that should apply during execution
- P produce “parseable” output for some commands that read registers
- i **<network-interface>**
the network interface (e.g. “en0”) to use for the command
- x print numeric values in hex rather than decimal

The default value for the -d **<device>** option can be set via the **OCPI_DEFAULT_HDL_DEVICE** environment variable.

The commands provided by the **ocpihdl** utility are summarized in the following table. The check-boxes show which ones require a device to be specified, a worker to be specified, or an Ethernet interface to be specified.

Table 22: The `ocpihdl` Commands

Name	Device	Worker	Interface	Description
<code>admin</code>	✓			Dump HDL device admin information
<code>bram</code>				Create BRAM file from input (XML) file
<code>deltatime</code>	✓			Measure round-trip time for synchronization
<code>emulate</code>			✓	Emulate a UDP or Ethernet device
<code>ethers</code>				List available/up/connected Ethernet interfaces
<code>probe</code>	✓			See if a specific device exists and responds
<code>load</code>	✓			Load a bitstream onto the device
<code>getxml</code>	✓			Extract XML metadata from running device
<code>radmin</code>	✓			Read addresses in HDL device admin space
<code>reset</code>	✓			Reset device (via control plane access)
<code>rmeta</code>	✓			Read metadata space from device
<code>search</code>			✓	List all discovered and responding HDL devices
<code>settime</code>	✓			Set the device's GPS time to the system time
<code>simulate</code>				Run a simulator server
<code>unbram</code>				Create an XML file from a BRAM file
<code>wadmin</code>	✓			Write admin space
<code>wclear</code>	✓	✓		Clear worker status errors
<code>wdump</code>	✓	✓		Print worker status registers (not properties)
<code>wop</code>	✓	✓		Perform control operation on worker (e.g. start)
<code>wread</code>	✓	✓		Read worker configuration property space
<code>wreset</code>	✓	✓		Assert reset for worker
<code>wunreset</code>	✓	✓		De-assert reset for worker
<code>wwctl</code>	✓	✓		Write worker control registers
<code>wwpage</code>	✓	✓		Write page register (to reach full 32 bit space)
<code>wwrite</code>	✓	✓		Write worker configuration property space

9.1 General, non-Worker Commands for the *ocpihdl* Utility

9.1.1 *admin* command – Print the Device’s Administrative Information

This command is used to dump all of the information and state in the devices “admin” space, which is the information for the device and loaded bitstream as a whole. A device flag must be specified. There are no command arguments.

9.1.2 *bram* command – Create a Configuration BRAM File from an XML File

This command converts an XML file that is expected to contain the “artifact description XML” into a format that will be processed into a read-only BRAM inside the FPGA bitstream. This allows software to know what is inside a bitstream by only looking at the device, without needing a separate bitstream file. This command is used by the scripts that create bitstreams. The output format is an ASCII format that the tools can use to initialize memories. The two command arguments are <input-file> and <output-file>, e.g.:

```
ocpihdl bram my.xml mybram.bin
```

9.1.3 *deltatime* command – Perform Time Synchronization Test on Device

This command, which requires a device option, uses special time-difference hardware in the OpenCPI FPGA bitstream to measure the round-trip time for accessing the FPGA in order to reduce the time-skew between the system time-of-day and the FPGA’s time-of-day. It takes 100 samples, averages them, eliminates 10% outliers, and then retests the time-skew after applying the correction.

9.1.4 *emulate* command – Emulate a Network-based Device (admin space only)

This command acts as an Ethernet or UDP-based HDL device and responds to discovery and admin-space accesses. It is used to test software controls and network connectivity. It can optionally be supplied with a network interface option (-i) to specify on which network interface should the emulated device appear. With no network interface specified, it uses the first available that it finds (that is “up” and “connected”). A special case is the network interface whose name is “udp”, which means the command emulates an OpenCPI HDL device attached to the IP subnet of the host computer implementing discovery and control via UDP, e.g.:

```
ocpihdl -i udp emulate
```

9.1.5 *ethers* command – Display Available (Up and Connected) Network Interfaces

This command is used to list all of the available Ethernet network interfaces on the system and whether they are “up” and “connected”. It also shows what the “default” interface is (the first listed that is up and connected) for commands that could take an interface (-i) option. It also shows the identity chosen for the system, based on the first interface with an address and a MAC address.

9.1.6 *probe* command – Test Existence and Availability of Device

The probe command takes a device (as an option or argument) and tries to contact it and see if it is responding. This should work whether it is running an application or not.

9.1.7 *load* command – Load Bitstream

The load command takes a device (as an option or argument) and a bitstream file name argument and loads the bitstream onto the device.

9.1.8 *getxml* command – Retrieve XML Metadata from Device

The getxml command takes a device (as an option or argument) and a filename to write the XML data to. It retrieves and uncompresses the XML data stored in the running device and writes it to a file.

9.1.9 *radmin* command – Read a Specific Address in Device’s Admin Space

This command reads an individual word of data from the admin space of the specified device. The command argument is the hexadecimal (starting with 0x) or decimal offset in the admin space. If the offset ends in “/n” where n is 4 or 8, then that specifies the size of the access in bytes. If there is no “/n”, then the access is a 32 bit access.

If the parseable option (-P) is specified the output is just the value returned in hexadecimal format. Otherwise a prettier message with the offset and value is printed. An example that does a 2 byte read of offset 12 would be:

```
ocpihdl -d pci:5 radmin 12/2
```

9.1.10 *reset* command – Perform Soft Reset on Device

This reset command resets the device in a way that does not affect the control path to the device. For example on a PCI-based device, it would not damage or reconfigure the PCI Express interface. All application workers and most other (infrastructure) workers are placed in reset and need to be specifically taken out of reset [unimplemented].

9.1.11 *rmeta* command – Read from Addresses in the Metadata Space of the Device

This command works exactly like the **radmin** command except it reads the configuration BRAM space rather than the admin space.

9.1.12 *search* command – Search for all Available HDL Devices

This command searches for all reachable HDL devices and reports what it finds. It uses all of the supported control paths (PCI, Ether, LSim). If an interface (-i) option is specified it limits the Ether search to that one interface.

9.1.13 *unbram* command – Create an XML File from a Config BRAM File

This command reverses the function of the **bram** command, by converting a file formatted for initializing (during assembly build) the configuration bram in a bitstream back to the original XML file. The two command arguments are <*input-file*> and <*output-file*>, e.g.:

```
ocpihdl unbram mybram.bin my.xml
```

9.1.14 **wadmin** command – Write Specific Addresses in the Device’s Admin Space

This debug command writes a 4-byte/32-bit or 8-byte/64-bit value in the device’s admin space at the specified offset. The size in bytes (4 or 8) is optionally specified in the first argument with a slash. The default size is 4. The syntax is:

```
ocpihdl wadmin -d <hdl-device> <offset>[/<size-in-bytes> <value>]
```

The value can be in hex (preceded by 0x) or decimal. An example, to write an 8 byte value at offset 0x20 with the value 12345, would be:

```
ocpihdl wadmin -d PCI:5 0x20/8 12345
```

9.1.15 **settime** command – Set Device’s Time from System Time

This command sets the time on the device from the current system time. It requires a device to be specified. The current time of the device is shown in the output of the **admin** command.

9.2 Worker Commands: Commands that Operate on Individual Workers

There are two sets of worker commands. The first set relies on embedded metadata and is generally more friendly and useful. The second set does not rely on metadata and is more primitive, and all of the latter begin with the letter **w**.

The following four commands take a worker instance name as the first argument, which is the same as the instance name in the HDL assembly. It can also be the index of the worker instance that is shown in previous commands, such as **ocpihdl get**, with no arguments.

9.2.1 *get command — Get All or Single Worker Instance Property Information*

With no arguments, this command displays information for all workers. With an argument to specify one instance, it displays information for that instance.

With the verbose option (**-v**), it will display all property values. Without the verbose option it will display only summary information. Using the hex (**-x**) option, all the numeric values will be printed in hexadecimal, and without it the values will be decimal.

When a worker is specified, an additional argument can specify the name of a property to display. Thus to display the prop1 value of the instX worker instance, in hex:

```
ocpihdl get -x instX prop1
```

This command knows the data types of properties and displays the values accordingly.

9.2.2 *set command — Set Property Value in a Worker Instance*

This command sets a particular property value of a particular worker instance. The second argument is the property name and the third argument is the value to set.

The syntax of the value is the same syntax used to specify an initial or default value of properties in OCS, OWD, or HDL assembly XML.

To set the value of a property that was an array of three shorts, an example would be:

```
ocpihdl set instX prop1 "-1,0x12,0177"
```

9.2.3 *control command — Change Control State of Worker Instance*

This command performs a control operation on the worker instance specified in the first argument. In addition to the defined life cycle control operations (initialize, start, stop, release, etc.), the operation can be **reset** and **unreset**. The operation is the second argument after the argument that specifies the worker instance. To put a worker in a reset condition (to assert reset), an example is:

```
ocpihdl control instX reset
```

9.2.4 *status command — Display the Status of a Worker Instance*

This command prints the status of the identified worker instance. In this context the status includes the control state, as well as the display output of the wdump command described below.

9.2.5 Primitive worker commands using worker indices

The following worker commands operate on either a single worker or a sequential range of workers. They all require a device to be specified, and the first argument is either a single worker number or a set of worker numbers separated by commas. Thus to perform a worker command on worker 5 would be:

```
ocpihdl w<cmd> -d pci:5 5 <worker command args if any>
```

To perform the command on workers 2, 5, and 6, would be:

```
ocpihdl w<cmd> -d pci:5 2,5,6 <worker command args if any>
```

The worker commands are:

9.2.6 **wclear** command – Clear a Worker’s Error Registers

Each worker has a set of control and status registers that are part of the control-plane infrastructure IP (not in the worker itself). This command clears the error and attention bits in the worker’s status register. A worker’s status can be displayed by the “wdump” command. An example for clearing error and attention status for worker 5 is:

```
ocpihdl wclear -d PCI:5 5
```

9.2.7 **wdump** command – Dump a Worker’s Status Registers (not Properties)

Each worker has a set of control and status registers that are part of the control-plane infrastructure IP (not in the worker itself). This command displays the current status of the worker by dumping these registers.

9.2.8 **wop** command – Perform a Control Operation on a Worker

This worker command executes a control operation for the worker, directly accessing the hardware that makes the control operation request of the worker. Only the “start” operation is implemented on all HDL workers. Any others may have unpredictable or erroneous results when requested on a worker that doesn’t implement them. The available operations are:

- **initialize** – after reset is deasserted (see wunreset), request that a worker initialize itself
- **start** – put the worker into an operational state, after stop or initialize
- **stop** – suspend operation of the worker
- **release** – return the worker to the “pre-initialized” state
- **test** – run the worker’s built-in test
- **before** – inform worker that a batch of property settings will happen
- **after** – inform a worker that a batch of property reads has completed

9.2.9 **wread** command – Read a Worker’s Configuration Property Space

This command reads a value from the worker’s property space. The first argument is the offset in the worker’s property space (in bytes), with an optional size-in-bytes for the

access (1, 2, 4 or 8), and the second (optional) argument is how many sequential accessed to make. The default for the size is 4, and the default for the second argument is 1. E.g., to read 3 single bytes at offset 6 from worker 11 would be:

```
ocpihdl read -d 5 11 6/1 3
```

See the “**wwpage**” command for workers whose property space is larger than 1 MByte (2^{20}).

9.2.10 **wreset** command – Assert Reset for a Worker

This command asserts the control reset signal into the worker. It stays asserted until the **wunreset** command is used.

9.2.11 **wunreset** command – Deassert Reset for a Worker

This command deasserts the control reset signal into the worker, after which the “wop initialize” command can be issued (or, if the worker does not implement the “initialize” command, the “wop start” command can be issued).

9.2.12 **wwctl** command – Write a Worker’s Control Register

This command writes the worker’s control register, which is in the control plane infrastructure IP, not in the worker itself. The argument is a 32 bit value. The bit definitions are described in the “OpenCPI HDL Infrastructure” document.

9.2.13 **wwpage** command – Write a Worker’s Window/Page Register

A worker’s property space can be a full 32 bit space (4 GBytes), but to access more than the first 1MByte (2^{20} bytes), a “window” register in the control-plane infrastructure must be set. This command sets that register with the value in the first argument, which sets the high order 12 address bits (31:20) of the effective address when the **wread** and **wwrite** commands are issued. The offset in those commands supplies bits (19:0) of the effective address. Thus to read location 0x12345678 in worker 7 with a full 4GByte property space, two commands would be used:

```
ocpihdl wpage -d pci:5 7 0x123  
ocpihdl wread -d pci:5 7 0x45678
```

9.2.14 **wwrite** command – Write a Worker’s Configuration Property Space

This command writes a single value into a worker’s property space at an offset (and size) specified in the first argument and a value specified in the second argument. The **wwpage** command applies to this command also for workers with large property spaces. Thus to write location 0x20 in worker 6 with the 64-bit value 0x123456789abc:

```
ocpihdl wwrite -d pci:5 6 0x20 0x123456789abc
```

10 HDL Platform and Device Development

HDL Platform development is the activity that makes a FPGA-based hardware platform fully enabled for running OpenCPI applications. This activity is sometimes called making a **Board Support Package (BSP)**. In OpenCPI the definition of an HDL platform is an FPGA surrounded by and connected to a set of devices, and possible slots that accept plug-in optional boards with devices on them.

Developing HDL platform workers and device workers is described in a separate ***OpenCPI Platform Development Guide*** document. It covers the development of HDL:

- Platform workers: the singleton worker that bootstraps the platform and container
- Device workers: workers that support external devices attached to FPGAs
- Platform configurations: assemblies of platform workers and some device workers
- Slot types: standard definitions of the signals and pins of a slot connector
- Slots on platforms: how to define slots on HDL Platforms
- Cards for slots: how to define cards that contain devices and plug into slots

All of these asset types can be developed in projects along with other assets.